

# Functional Concurrency with F#

BY RICCARDO TERRELL - @TRIKACE  
[TERICCARDO@GMAIL.COM](mailto:TERICCARDO@GMAIL.COM)

# Agenda

Message Passing programming model & motivation

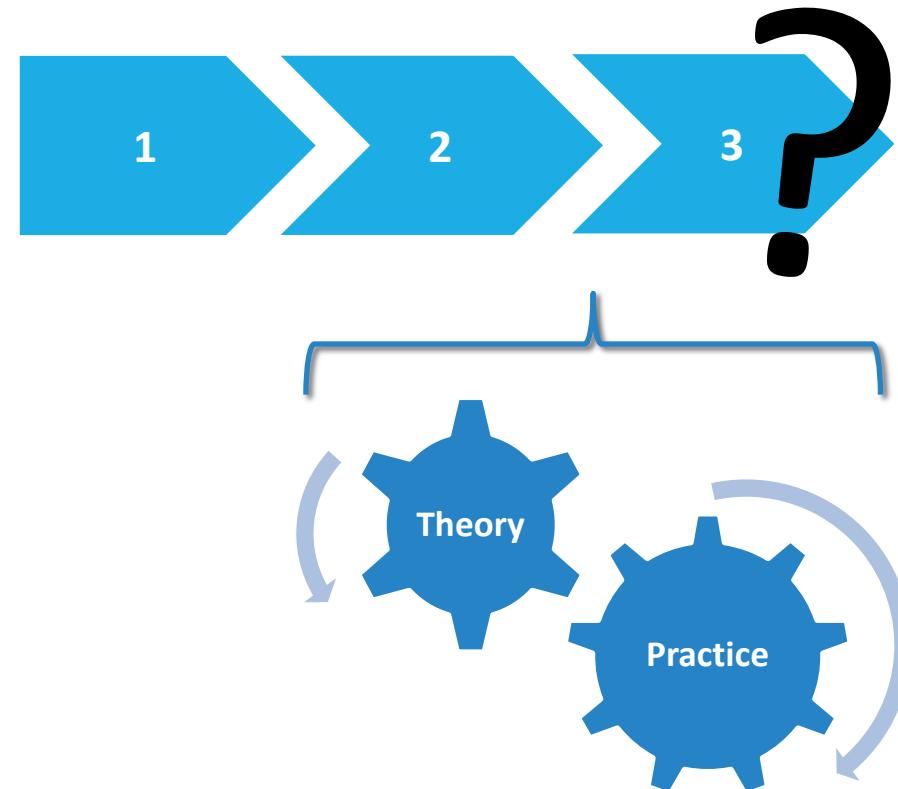
How to parallelize Actors/Agents seamlessly

Combinators to compose Asynchronous operations and Agent/Actors

Message Passing and Reactive programming

# Modules

---



# Introduction - Riccardo Terrell

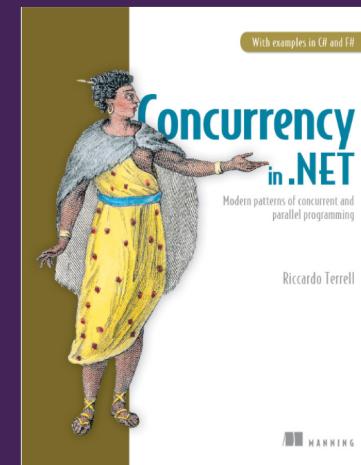
- ④ Originally from Italy, currently - Living/working in Washington DC
- ④ +/- 20 years in professional programming
  - ④ C++/VB → Java → .Net C# → Scala → Haskell → C# & F# → ??
- ④ Author of the book “Functional Concurrency in .NET” – Manning Pub.
- ④ *Polyglot programmer - believes in the art of finding the right tool for the job*
- ④ *Organizer of the DC F# User Group*



@trikace

[www.rickyterrell.com](http://www.rickyterrell.com)

[tericcardo@gmail.com](mailto:tericcardo@gmail.com)



<https://ti.to/luteceo-workshops/concurrent-functional-programming-in-net/en>



**LUTECEO**  
workshops

Concurrent Functional Programming in .NET

📅 April 25th–26th, 2019    📍 Paris, France

A photograph of Riccardo Terrell, a man with glasses and a dark sweater, speaking at a conference. He is gesturing with his right hand. To his left is a blue circular button with the word "Courses". Below the photo is a QR code and the "skills matter" logo.

**Functional  
Concurrency in  
.NET with C# & F#  
with Riccardo  
Terrell**

A two day course to master the necessary practices to build concurrent and scalable programs in .NET

Visit [skillsmatter.com/courses](http://skillsmatter.com/courses) for more information

**29th - 30th April 2019 in London**

<https://skillsmatter.com/courses/589-functional-concurrency-in-dotnet-with-csharp-and-fsharp-with-riccardo-terrell>



**What about you?**

Download links:

<https://github.com/rikace/funconcurrency-ws>

See **README** section pre-requisites

# Pre-requisites

---

## > Windows

- > Visual Studio Community or higher
- > Visual Studio Code
  - > C# Extensions
  - > F# Compiler + Ionide package
- > Rider

## > Linux

- > Visual Studio Code + (Mono or dotnetcore) + Ionide package

## > Mac

- > Visual Studio for Mac + (Mono or dotnetcore) or
- > Visual Studio Code + (Mono or dotnetcore) + Ionide package
- > Rider

Download links:

<https://github.com/rikace/funconcurrency-ws>

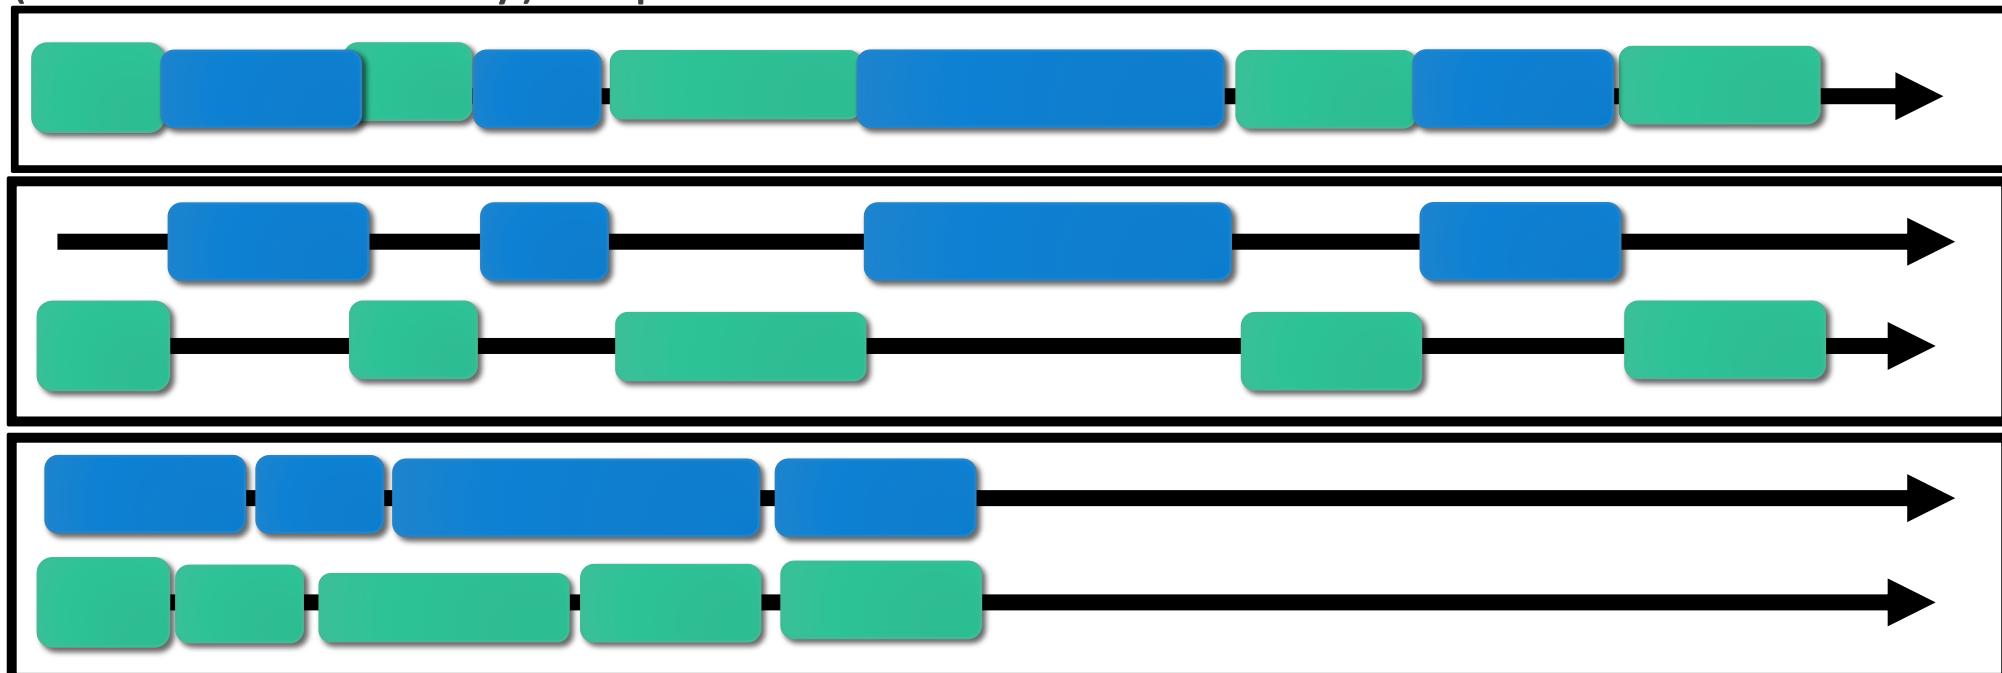
See README section pre-requisites

We need a way to make sure our programs continue to improve their speed and performance as hardware evolves. Have any ideas?



# Concurrency programming

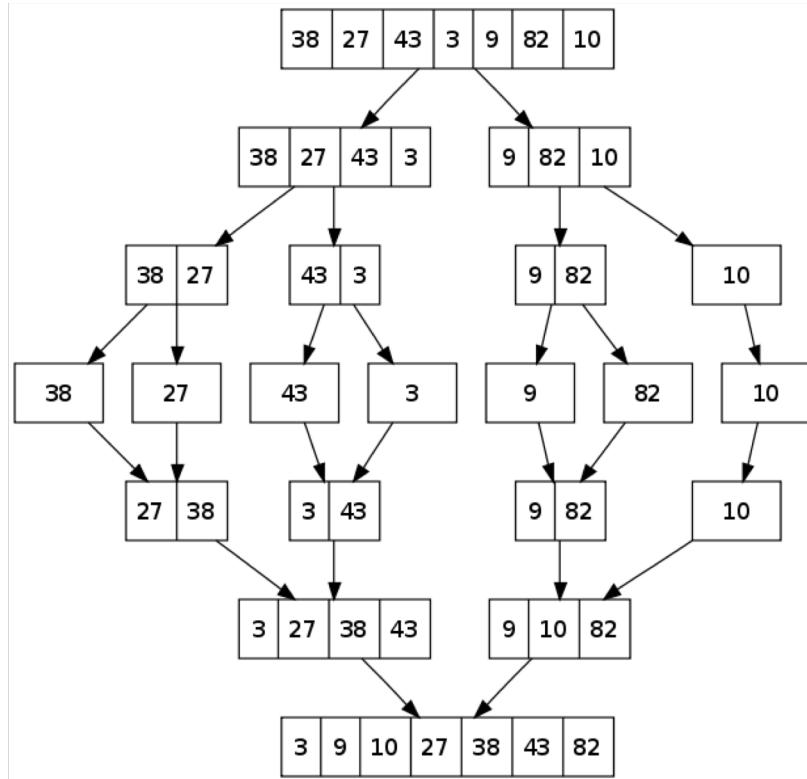
Concurrency provides a way to **structure a solution** to solve a problem that may (but not necessarily) be parallelizable.



Concurrency is the **composition** of independently executing computations.

# Divide and Conquer

---



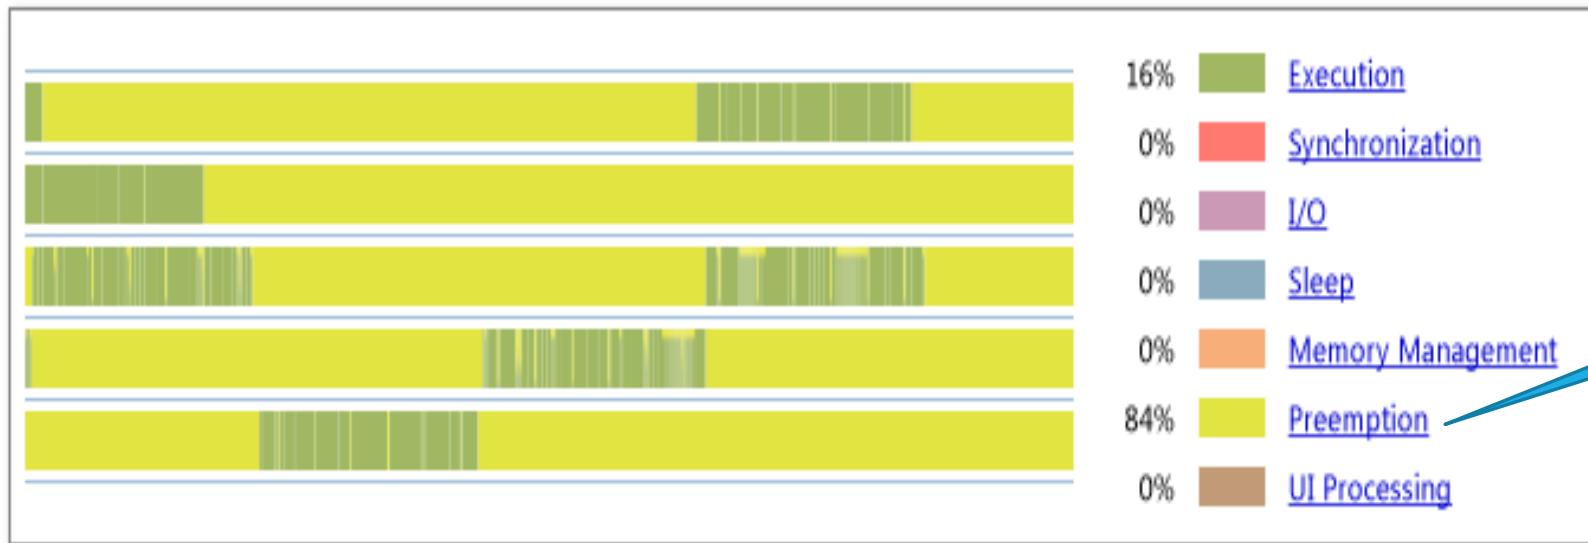
```
void QuickSort_Parallel<T>(T[] items, int left, int right)
{
    int pivot = left;

    SwapElements(items, left, pivot);

    Parallel.Invoke(
        () => QuickSort_Parallel(items, left, pivot - 1),
        () => QuickSort_Parallel(items, pivot + 1, right)
    );
}
```

# Quick Sort – Oversubscription

---

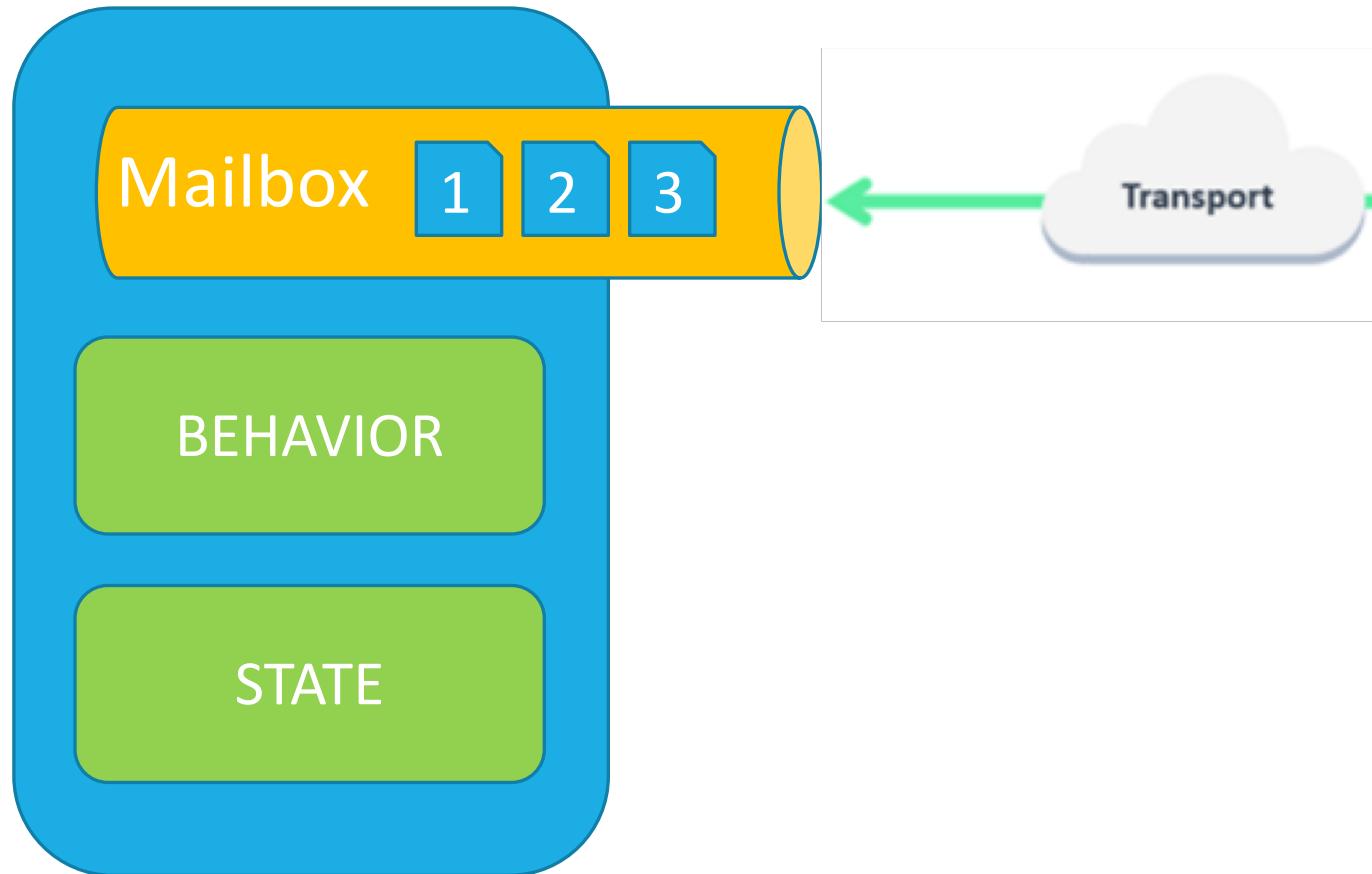


Preemption

# Message Passing

# Message Passing based concurrency

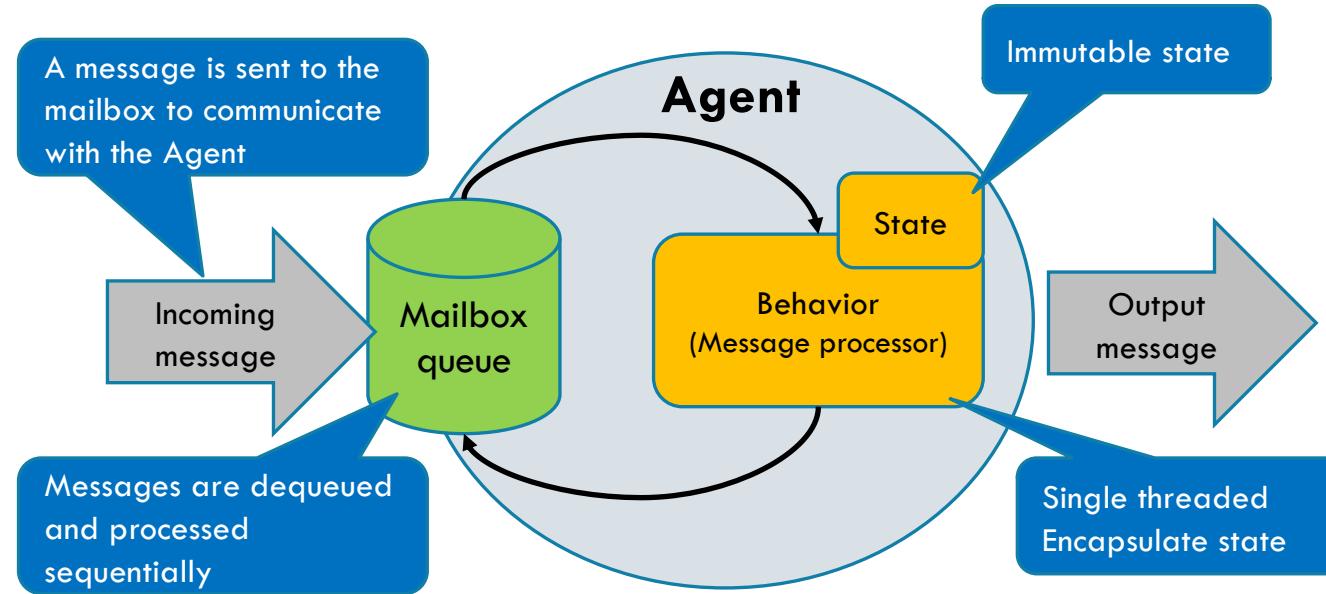
---



- Processing
- Storage – State
- Communication only by messages
- Share Nothing
- Message are passed by value
- Lightweight object
- Running on it's own thread
- No shared state
- Messages are kept in mailbox and processed in order
- Massively scalable and lightening fast because of the small call stack

# Agent anatomy

---

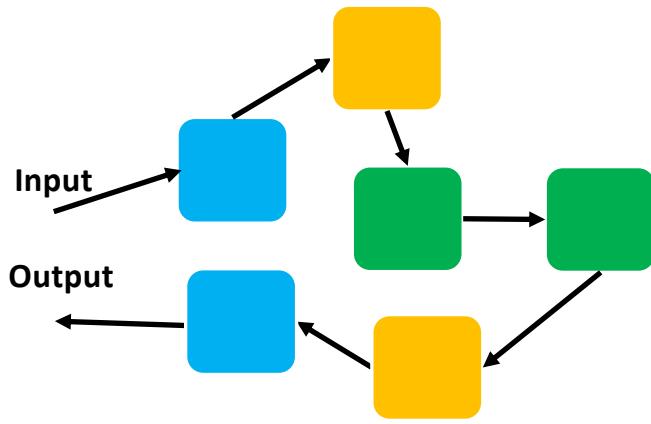


- Communication between agents done by **sending message**
- The messages are put on a **queue** for each agent
- Each agent process **one message** at time

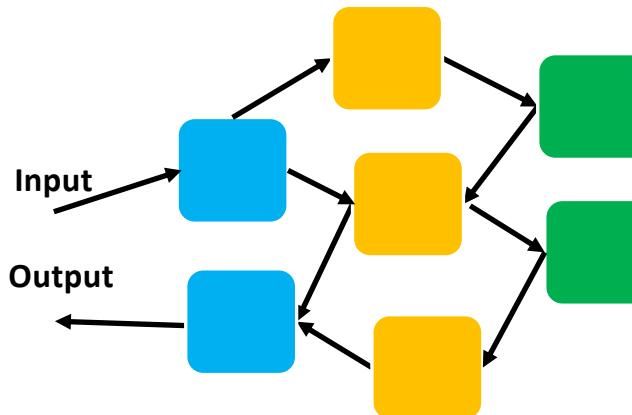
# Comparison between Sequential, Task-based and Message passing programming

---

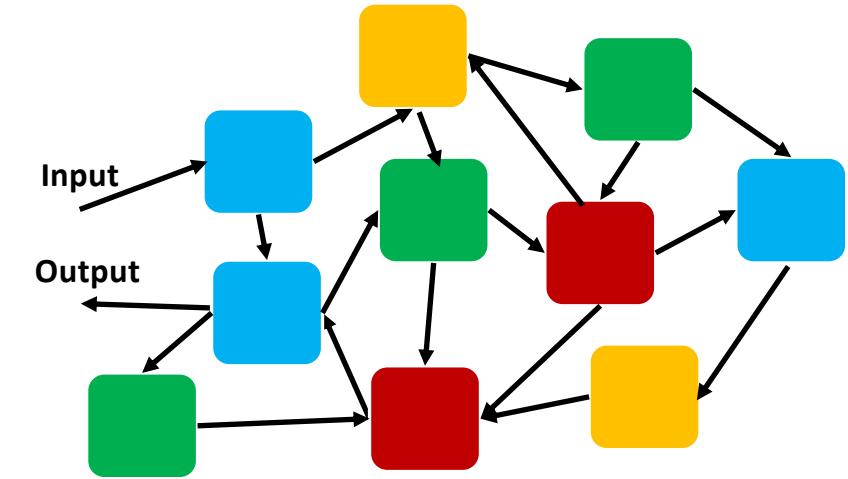
**Sequential Programming**



**Task-Based Programming**

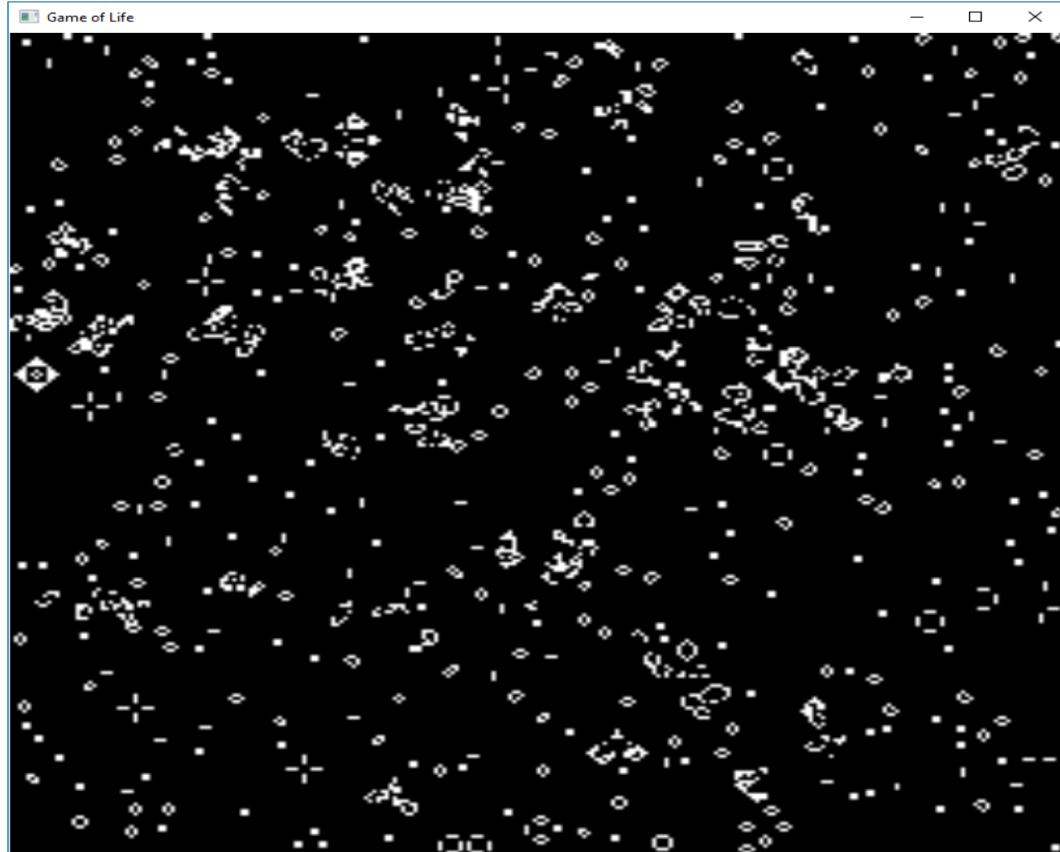


**Message-Passing Programming**



# Game of Life

---



The Game of Life rules:

- Each cell with one or no neighbors dies, as if by solitude.
- Each cell with four or more neighbors dies, as if by overpopulation.
- Each cell with two or three neighbors survives.
- Each cell with three neighbors becomes populated (maximum).

# Simple agent in F#

---

Receive message and say “Hello”

```
let hello = MailboxProcessor.Start(fun agent -> async {
    while true do
        let! name = agent.Receive()
        printfn "Hello %s" name
        do! Async.Sleep 1000 })

hello.Post("World!")
```

- Single instance of the body is running
- Waiting for message is asynchronous
- Messages are queued by the agent

# Mutable and immutable state

---

## Mutable state

- Accessed from the body
- Used in loops or recursion
- Mutable variables (ref)
- Fast mutable collections

## Immutable state

- Passed as an argument
- Using recursion (**return!**)
- Immutable types
- Can be returned from the agent

```
type Agent<'a> = MailboxProcessor<'a>
```

```
Agent.Start(fun agent -> async {
    let names = ResizeArray<_>()
    while true do
        let! name = agent.Receive()
        names.Add(name) })
```

```
Agent.Start(fun agent ->
    let rec loop names = async {
        let! name = agent.Receive()
        return! loop (name::names) }
    loop [])
```

# Declaring messages

---

Agents handle multiple messages

- Message type using discriminated union

```
type internal OnePlaceMessage<'T> =
| Put of 'T
| Get of AsyncReplyChannel<'T>
```

Safety guarantees

- Agent will be able to handle all messages

# Send message to agent

---

```
printerAgent.Post "hello"
```

```
printerAgent.Post "hello again"
```

```
printerAgent.Post "hello a third time"
```

# Encapsulating agents

---

Create type and expose messages as members

```
type internal OnePlaceMessage<'T> =
| Put of 'T
| Get of AsyncReplyChannel<'T>

type OnePlaceAgent<'T>() =
    let agent = Agent.Start(fun agent -> (* ... *))
    member x.Put(value) = agent.Post(Put value)
    member x.AsyncGet() = agent.PostAndAsyncReply(Get)
```

- Ordinary methods for encapsulating Post
- Asynchronous methods when waiting for a reply

# Agent Replying to the sender

---

**Message** carries input and a callback

```
type Message = string * AsyncReplyChannel<string>
```

**Reply** using the callback object

```
let echo = Agent<Message>.Start(fun agent ->
    async { while true do
        let! name, rchan = agent.Receive()
        rchan.Reply("Hello " + name) })
```

**Asynchronous** communication

```
let! s = echo.PostAndAsyncReply(fun ch -> "F#", ch)
```

# Agent State Machine

---

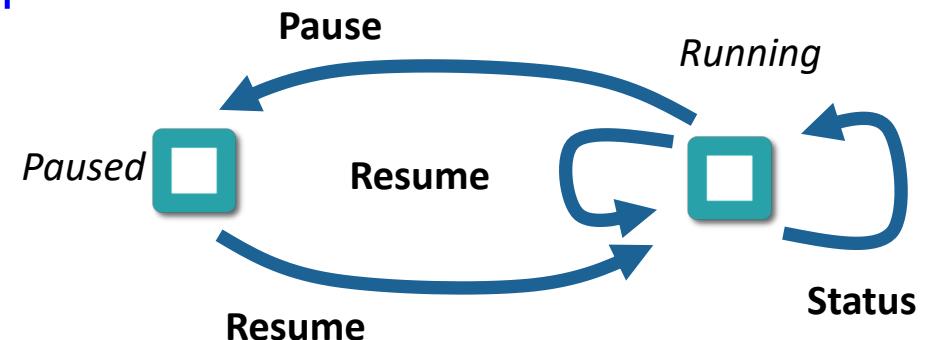
Multi-state agents use state machines

- Easy to implement as recursive functions
- Some states may leave messages in the queue

State transitions

- Accepting all messages, Asynchronously Receive and use pattern matching
- Waiting for a specific message, other messages stay in the queue

```
Agent.Start(fun agent ->
    let rec paused = agent.Scan (function
        | Resume -> Some (async {
            printfn "Resumed!"
            return! running })
        | _ -> None)
    and running = (* ... *) )
```

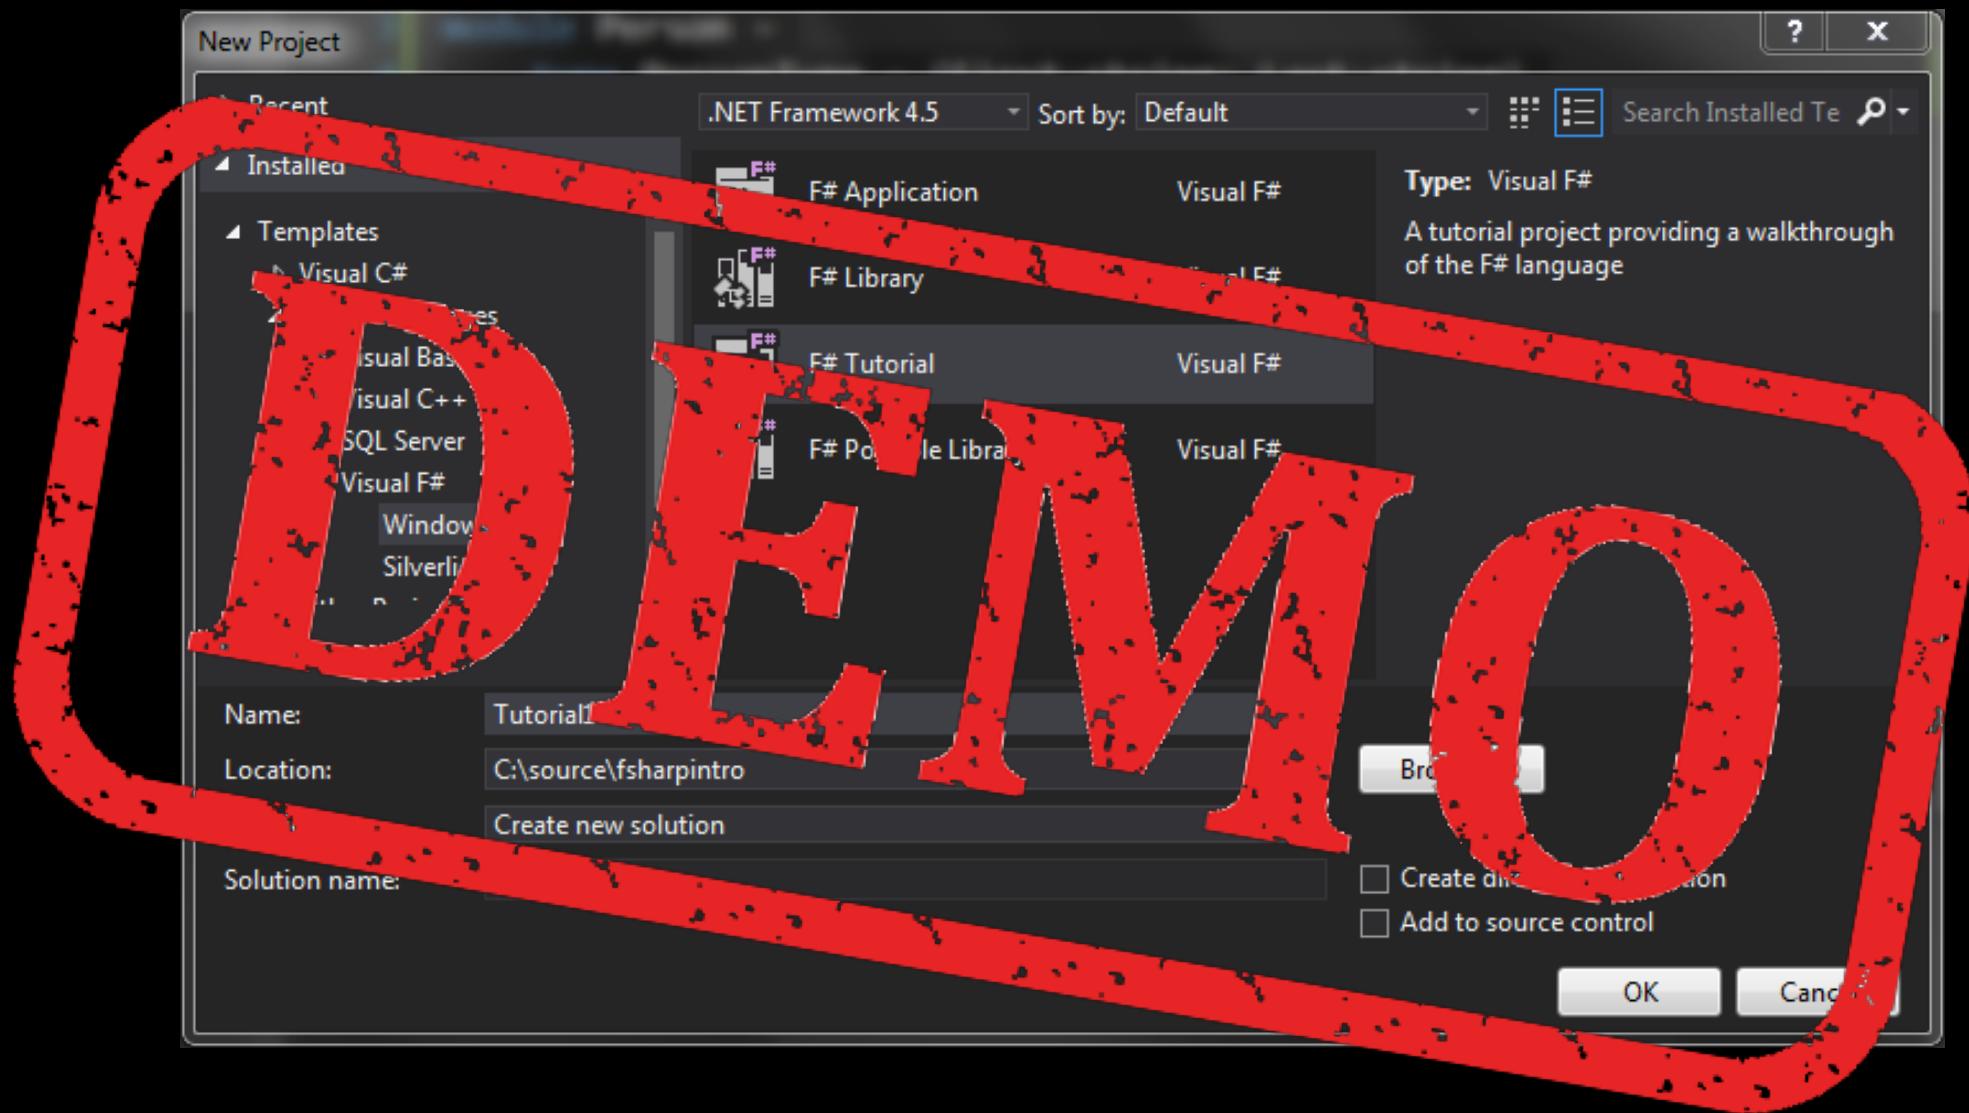


# Agent Error Handling & Disposable

---

```
let errorAgent =
    Agent<int * System.Exception>.Start(fun inbox ->
        async { while true do
            let! (agentId, err) = inbox.Receive()
            printfn "an error '%s' occurred in agent %d" err.Message agentId
            inbox.Post((agentId, err))
        }
    )

let agent cancellationToken =
    new Agent<string>((fun inbox ->
        async { while true do
            let! msg = inbox.Receive()
            failwith "fail!" }, cancellationToken.Token)
    )
    agent.Error.Add(fun error -> errorAgent.Post (error))
    agent.Start()
    agent
```



# Agent Patterns

# Agent-based architecture

---

Using loosely coupled connections

- Agents don't reference each other directly

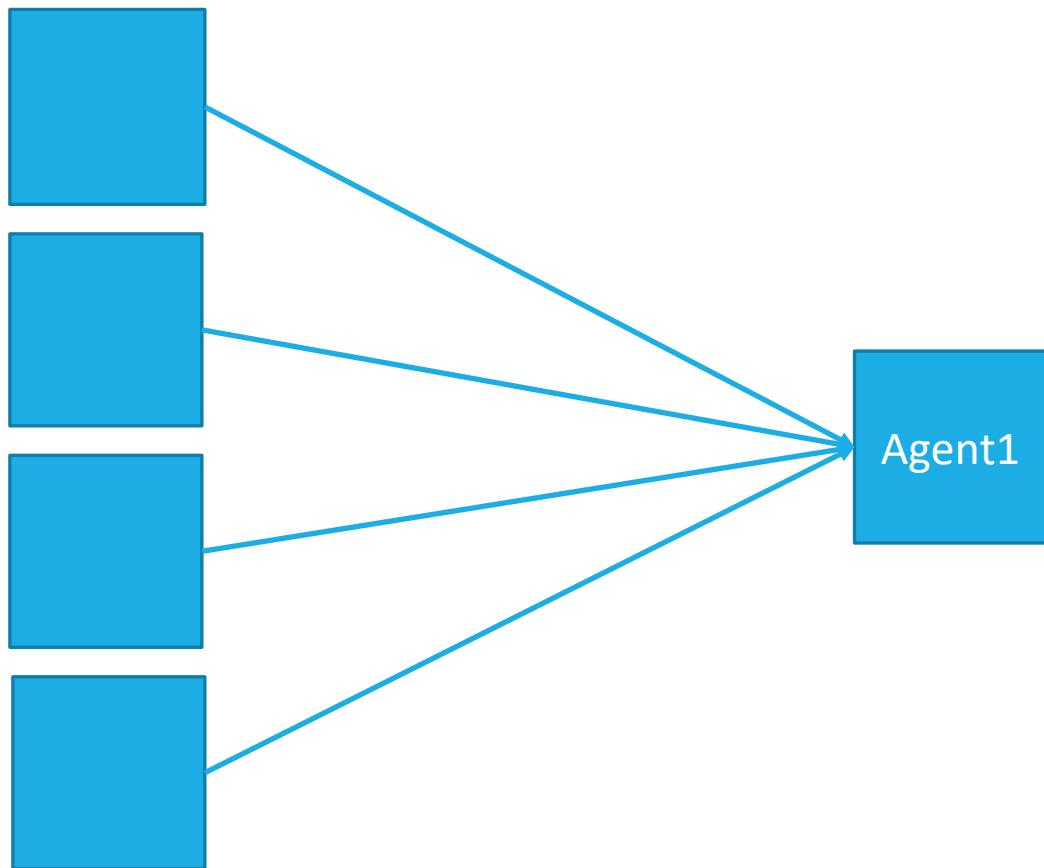
Common ways of organizing agents

- **Worker agent** – Single agent does work in background
- **Layered network** – Agent uses agents from lower level
- **Pipeline processing** – Step-by-step processing

# Agent patterns

---

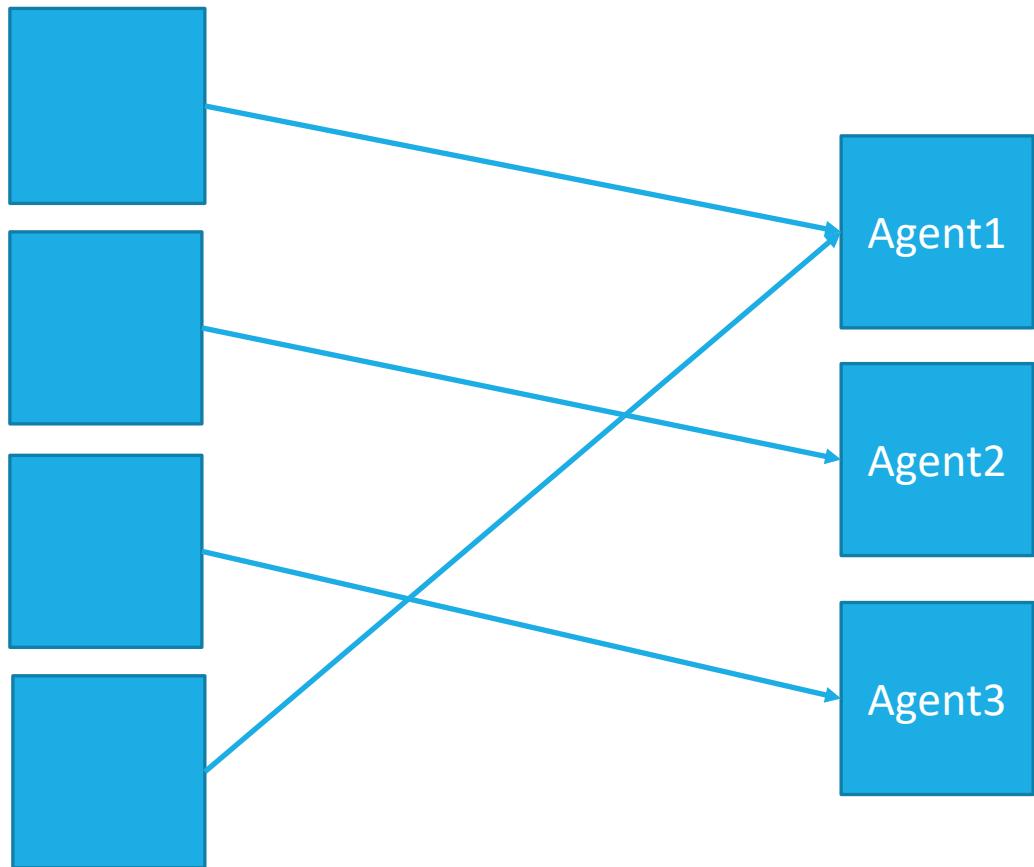
Singleton



# Agent patterns

---

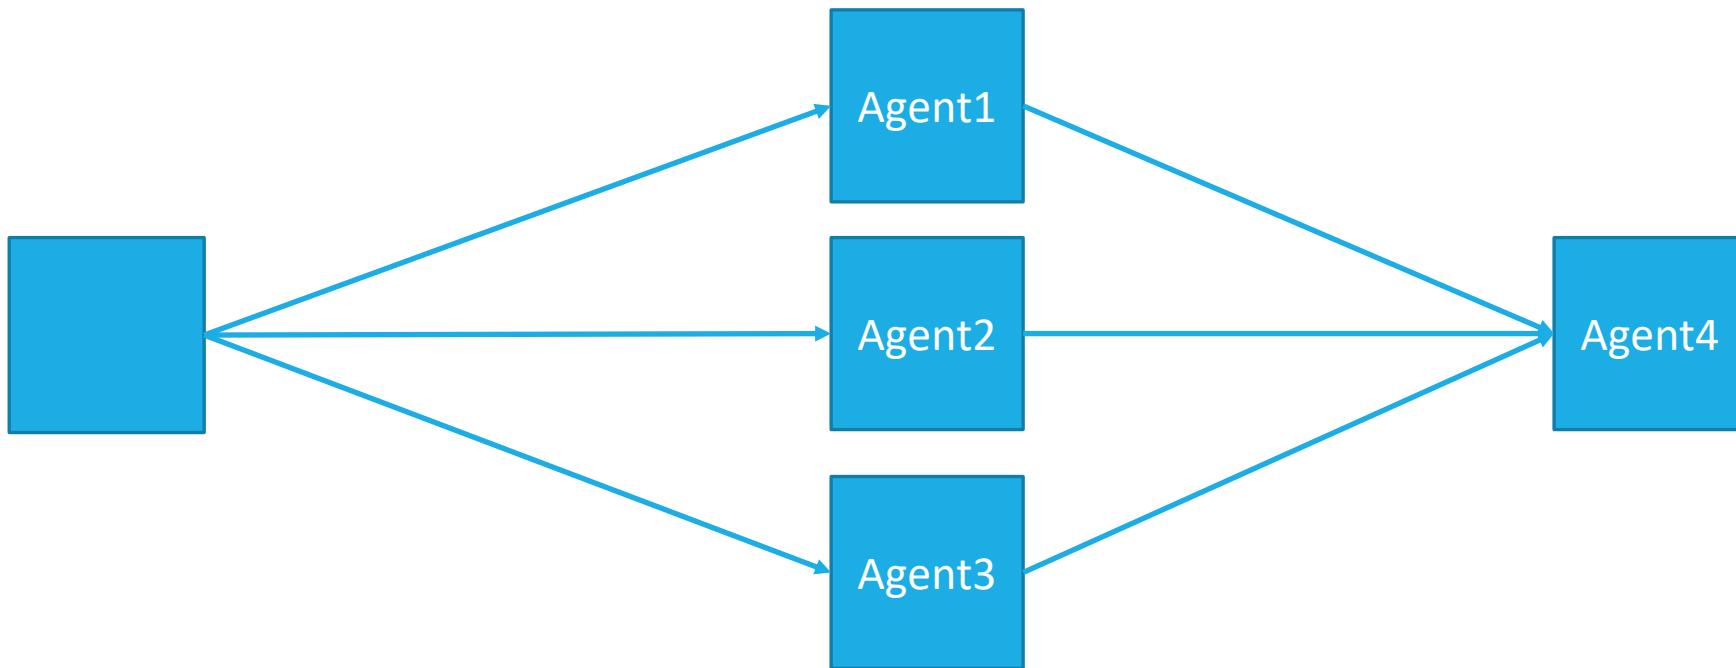
Pool



# Agent patterns

---

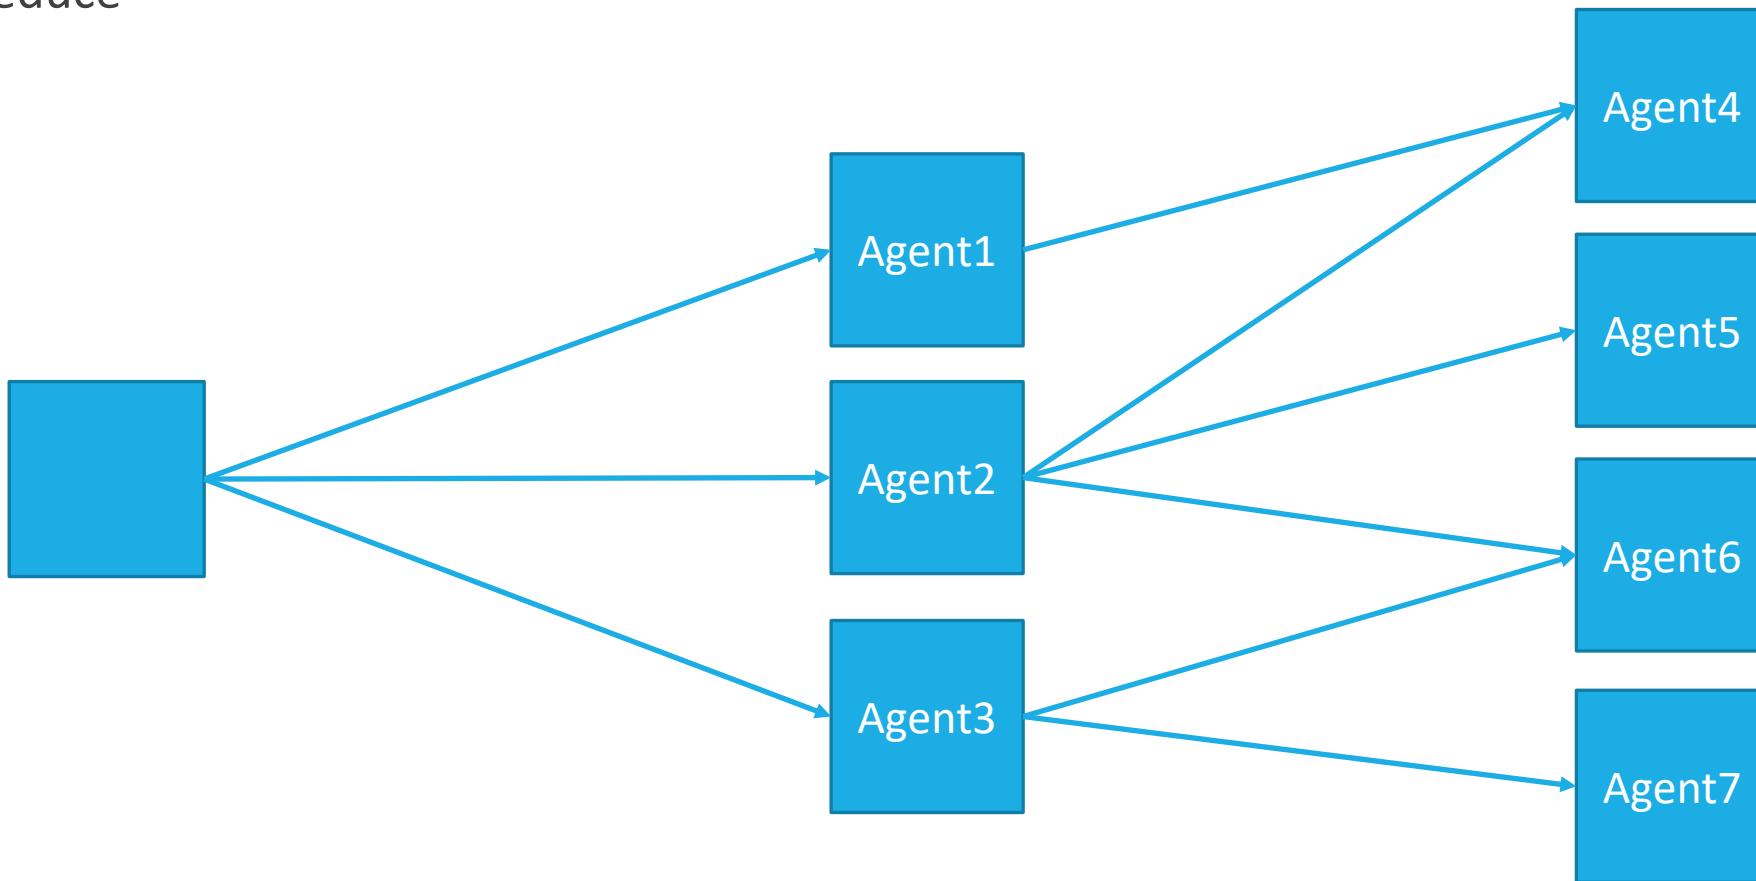
Fork Join



# Agent patterns

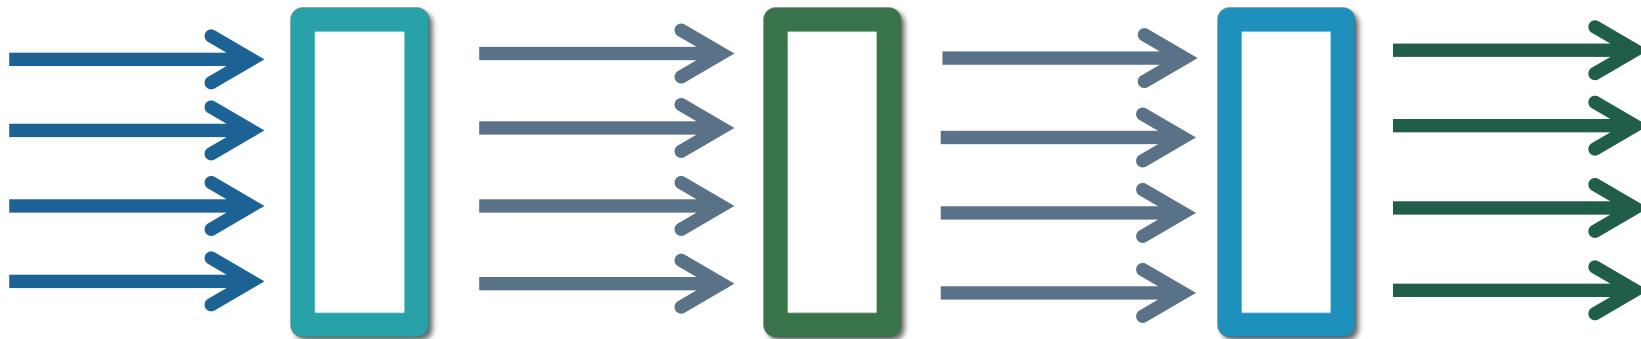
---

Map Reduce



# Pipeline Processing

---

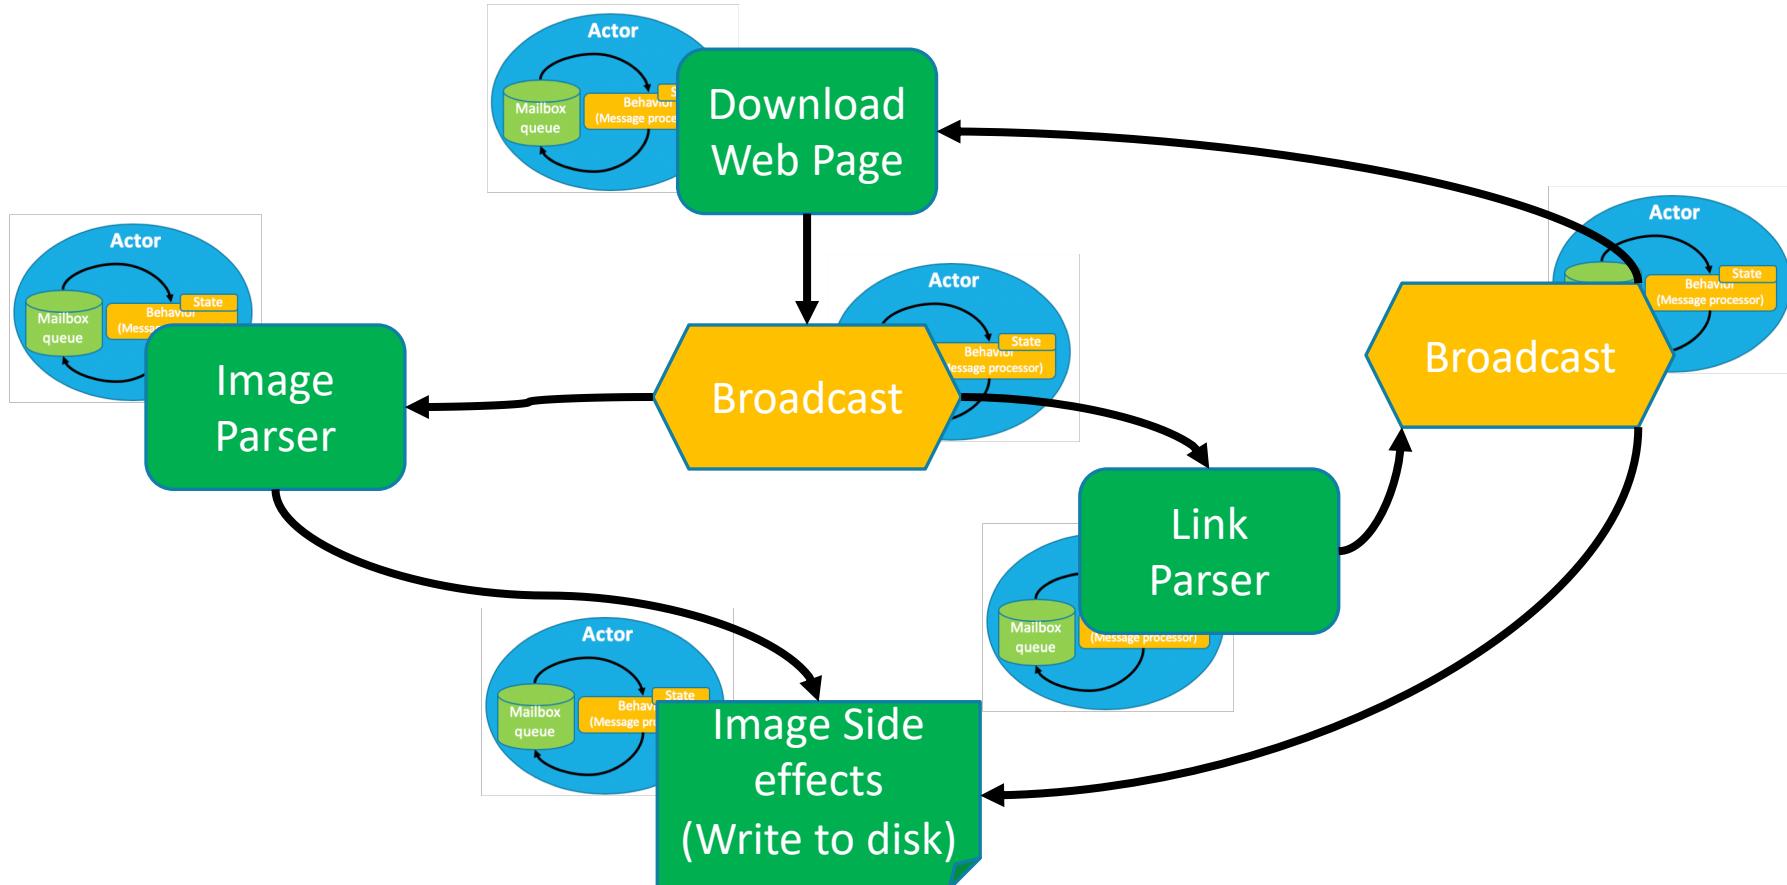


- A pipeline is a set of data processing elements connected in series, so that the output of one element is the input of the next one. The elements of a pipeline are often executed in parallel or in time-sliced fashion; in that case, some amount of buffer storage is often inserted between elements

# Parallel Web-Crawler

# Actors are hard to compose

## An Actor-based WebCrawler



# Web Crawler with Agents (assumptions)

---

- Web crawler must not spam websites, so each request to a single website should be delayed.
- All operations must be non-blocking.
- We assume that the web connection is not good
- Since we will be running the crawler on a regular computer, then we need to stop the system when we scrap enough links. The check itself will be regulated by a simple threshold value.

# Web-Crawler elements

---

- **Downloader:** the responsibility of the downloader is to download the html of a web page. It takes a url as input and it produces the related page's html as output.
- **Link-parser:** analyzes the downloaded html to look for links elements <a href="...">
- **Image-parser:** analyzes the downloaded html to look for image elements <image src="...">
- **Broadcast:** the broadcast block is construct from an input and output of single item linking one or more agents
- **Link-together (broadcast):** linking the block to each other
- **Writer:** persists the image and/or produce a side effect

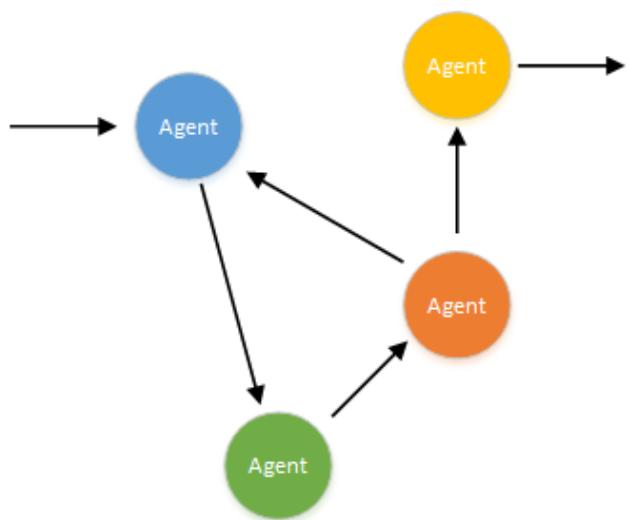
the **downloader** should be link to the **content broadcaster** which in tern should be linked both to the **image and link parser**, the **image parser** should be linked to the **writer** and the **link parser** should be linked back to the **downloader** (so it can crawl farther)

# Lab: Agent WebCrawler

# Agent Composition

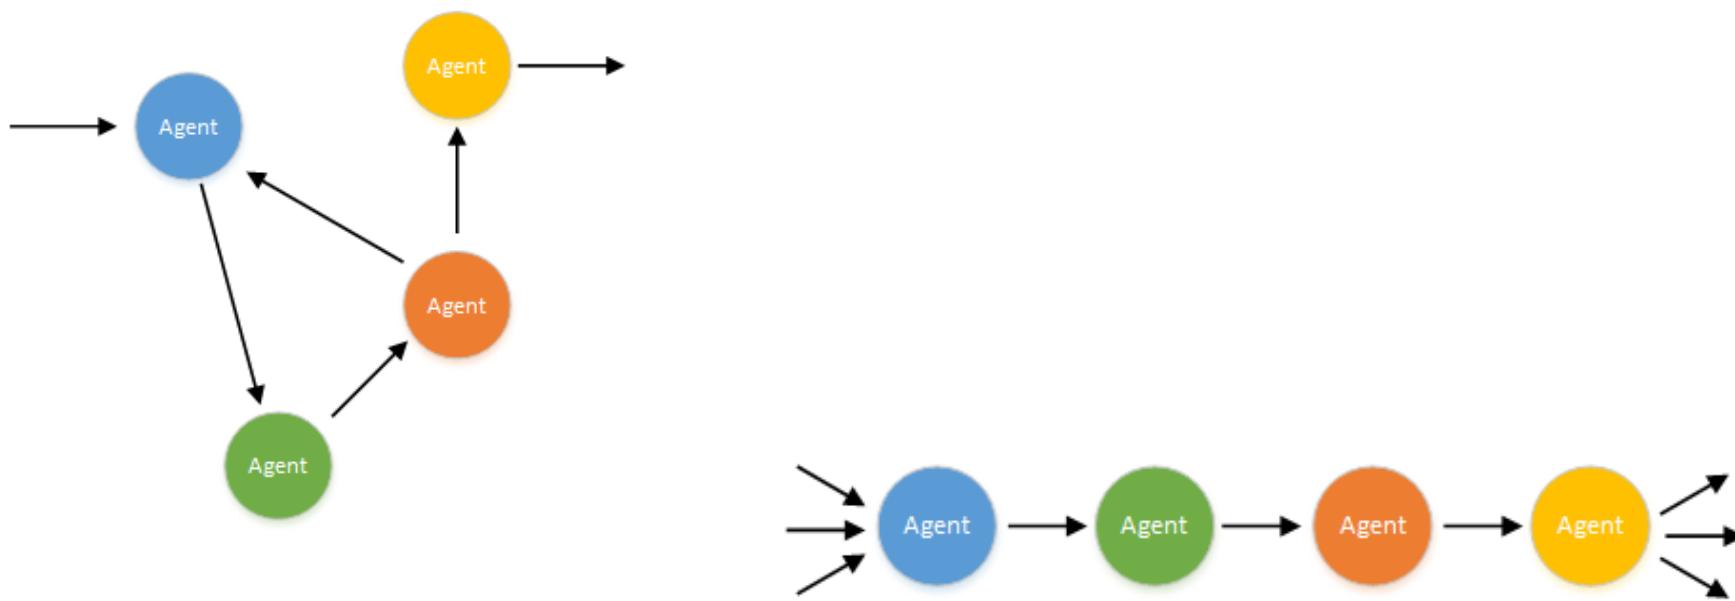
# Agent Pipeline

---



# Agent Pipeline

---



# Agent Pipeline

---

```
type Agent<'a> = MailboxProcessor<'a>
type Message = string * AsyncReplyChannel<string>

let agent f = Agent<Message>.Start(fun inbox ->
    let rec loop () = async {
        // TODO

        return! loop() }

    loop() )
```

# Agent Binding

---

**('a -> Async<'b>) -> Async<'a> -> Async<'b>**

```
let bind f xAsync = async {
    let! x = xAsync
    return! f x }
```

```
let retn x = async { return x }
```

```
let (>>=) xAsync f = bind f xAsync
```

# Agent Pipeline

---

```
let agent f = Agent<Message>.Start(fun inbox ->
    let rec loop () = async {
        let! msg, replyChannel = inbox.Receive()
        f msg
        replyChannel.Reply msg

        return! loop() }
    loop() )
```

# Agent Binding

---

```
let (>>=) x f = Async.bind f x  
let pipeline x = Async.retn x >>= agent1 >>= agent2
```

```
let (>=>) f1 f2 x = f1 x >>= f2  
let pipeline = agent1 >=> agent2
```

pipeline message

# Lab: Agent Composition

# Agent Pipeline

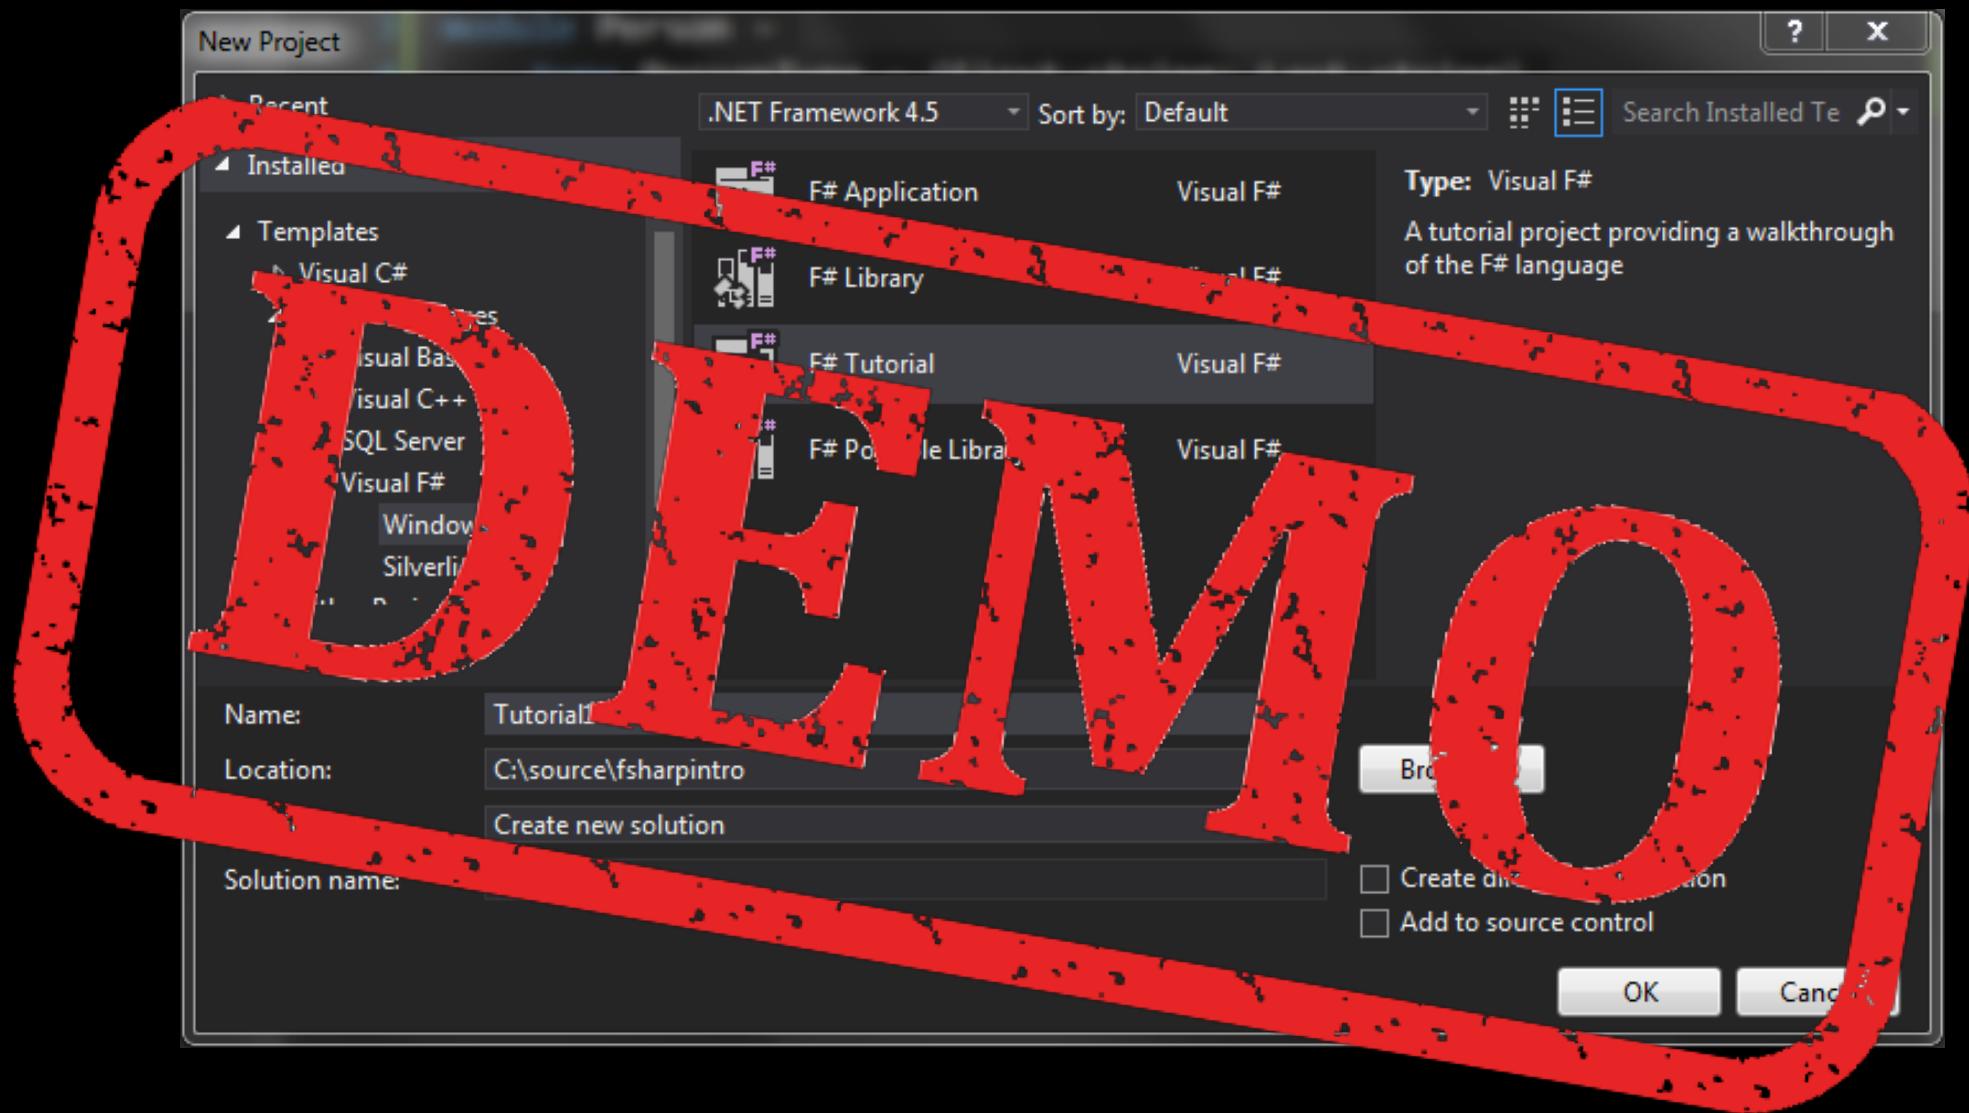
---

```
let agent f = Agent<Message>.Start(fun inbox ->
    let rec loop () = async {
        let! msg, replyChannel = inbox.Receive()
        f msg
        replyChannel.Reply msg

        return! loop() }
    loop() )

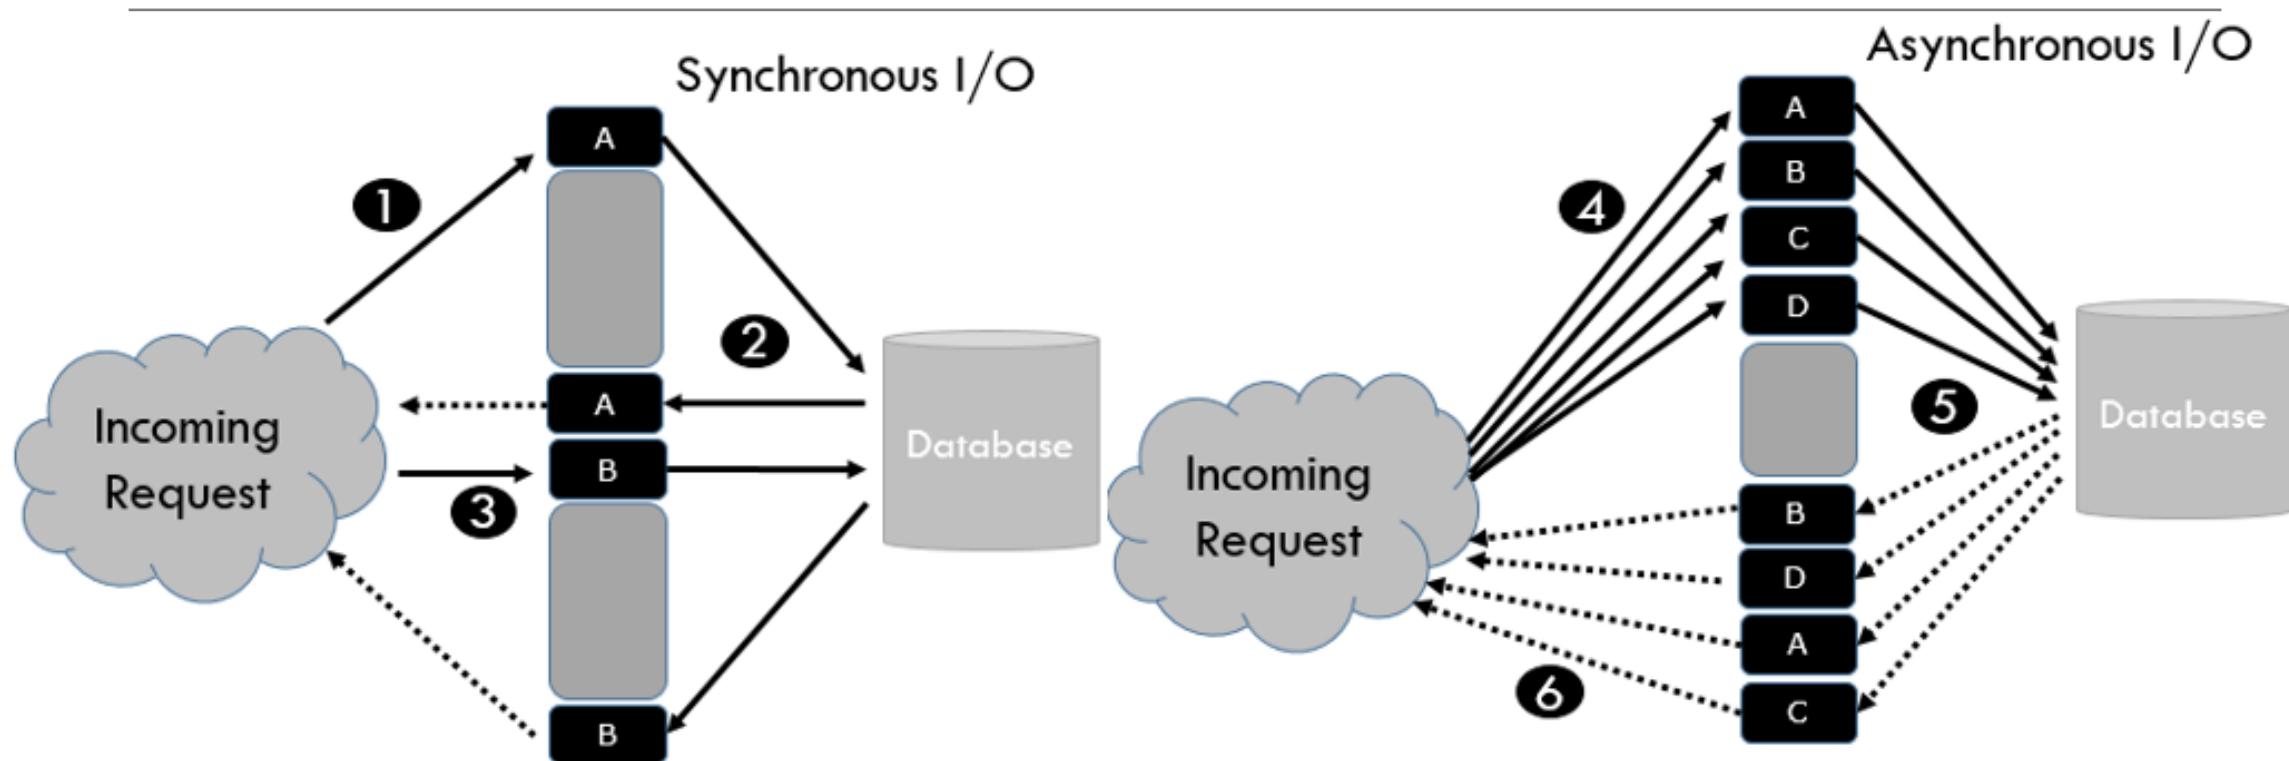
let pipelineAgent f m =
    let a = agent f
    a.PostAndAsyncReply(fun replyChannel -> m, replyChannel)

let agent1 = pipelineAgent (printfn "Pipeline processing 1: %s")
let agent2 = pipelineAgent (printfn "Pipeline processing 2: %s")
```



# Asynchronous Programming and Composition

# Synchronous vs Asynchronous



*“If you do nothing, you can scale infinitely.”*

*— Hanselman’s rule of scale*

# Its all about Scalability

---

```
let httpAsync (url : string) = async {
    let req = WebRequest.Create(url)
    let! resp = req.AsyncGetResponse()
    use stream = resp.GetResponseStream()
    use reader = new StreamReader(stream)
    return! reader.ReadToEndAsync() }

let sites =
    [ "http://www.live.com"; "http://www.fsharp.org";
      "http://news.live.com"; "http://www.digg.com";
      "http://www.yahoo.com"; "http://www.amazon.com"
      "http://www.google.com"; "http://www.netflix.com";
      "http://www.facebook.com"; "http://www.docs.google.com";
      "http://www.youtube.com"; "http://www.gmail.com";
      "http://www.reddit.com"; "http://www.twitter.com"; ]

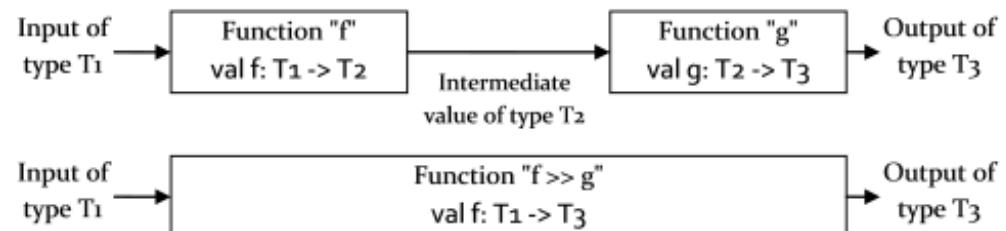
sites
|> Seq.map httpAsync
|> Async.Parallel
|> Async.Start
```

# Function Composition

---

**Function Composition - Building with composition.** **Composition is the 'glue'** that allows us build larger systems from smaller ones. This is the very heart of the functional style. Almost every line of code is a composable expression. Composition is used to build basic functions, and then functions that use those functions, and so on.

**Composition** — in the form of passed parameters plus first-class functions



# How can we compose Async?

```
Async<int> runOne () -> // ...
```

```
Async<decimal> runTwo(input : int) -> // ...
```

```
let result = runTwo(runOne()) // Error!!
```

F#

```
Async<int> runOne () -> // ...
```

```
Async<decimal> runTwo(input : int) -> // ...
```

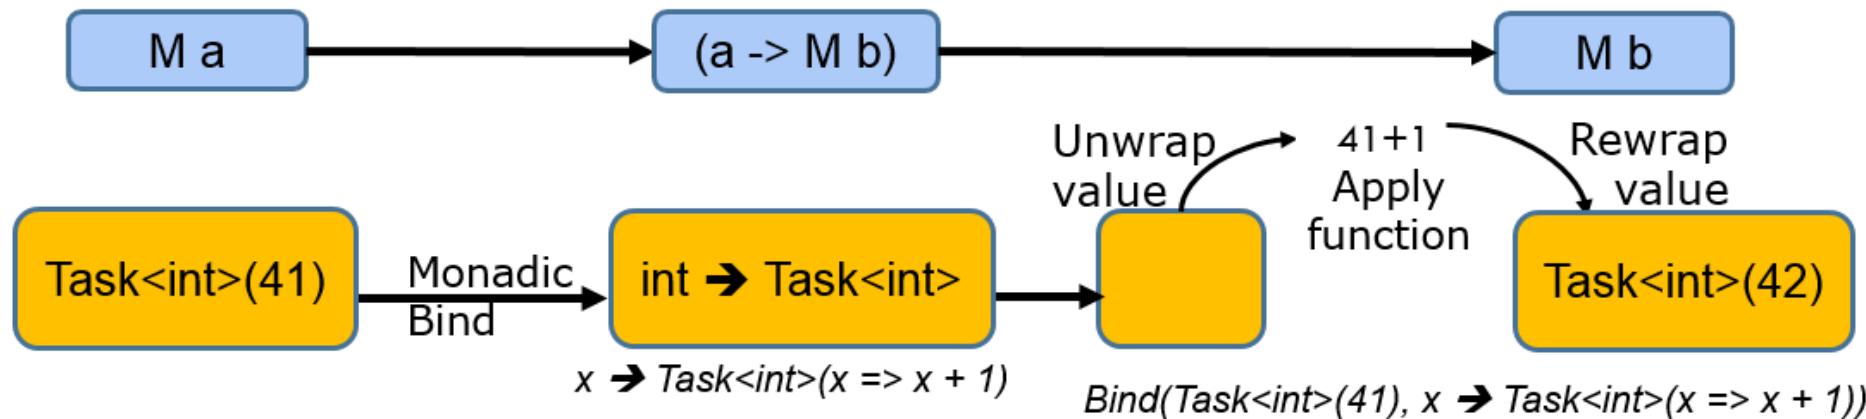
```
let compose (f1 : a -> Async<b>) (f2 : b -> Async<c>) =
    f1 >> f2 // Error!!
```

F#

# Composing Async (and Task)

Async<R> Bind<T, R>(Async<T> m, k : T -> Async<R>) ...

Async<T> Return(T value) ...

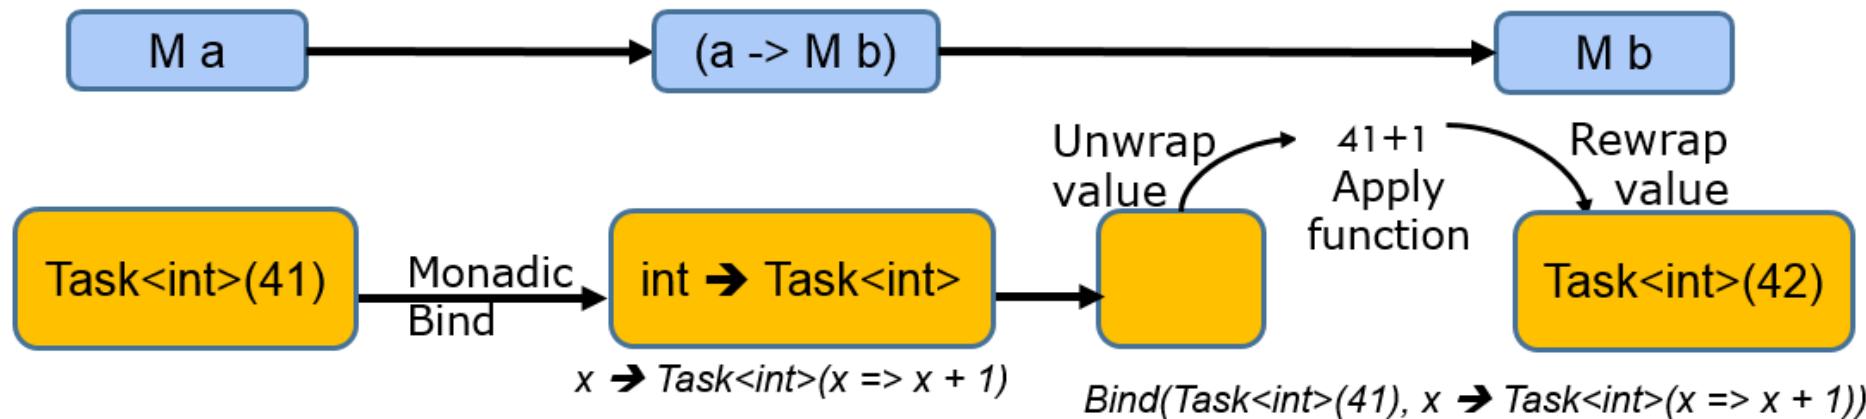


# Functional Design Patterns – Monad

Generic type

With function ( $>>=$ ): bind :  $M<'a> \rightarrow ('a \rightarrow M<'b>) \rightarrow M<'b>$

With function: return :  $'a \rightarrow M<'a>$



# Task/Async Bind and Return

---

```
// 'T -> M<'T>
static member Return (value : 'T) : Task<'T>

// M<'T> * ('T -> M<'U>) -> M<'U>
static member Bind (input : Task<'T>, binder : 'T -> Task<'U>) : Task<'U>

// 'T -> M<'T>
static member Return (value : 'T) : Async<'T>

// M<'T> * ('T -> M<'U>) -> M<'U>
static member Bind (input : Async<'T>, binder : 'T -> Async<'U>) : Async<'U>
```

# How can we compose Async?

```
Async<int> runOne () -> // ...
```

```
Async<decimal> runTwo(input : int) -> // ...
```

```
let compose = async {
    let! res = runOne()
    return! runTwo res
}
```

F#

# Kliesli operator in F#

---

```
fAsync: ('a -> Async<'b>) -> gAsync: ('b -> Async<'c>) -> arg:'a -> Async<'c>

let Kliesli (fAsync:'a -> Async<'b>) (gAsync:'b -> Async<'c>) (arg:'a) = ...

let (>=>) = Kliesli

let processStockHistory =
    asyncOne (_ -> Async<string>) >=> asyncTwo (string -> Async<int>)
```

# Lab: Async Composition

# Error handling in functional way

# Async Futures - Twitter Paper

---

## Your Server as a Function

Marius Eriksen

Twitter Inc.

marius@twitter.com

### Abstract

Building server software in a large-scale setting, where systems exhibit a high degree of concurrency and environmental variability, is a challenging task to even the most experienced programmer. Efficiency, safety, and robustness are paramount—goals which have traditionally conflicted with modularity, reusability, and flexibility.

We describe three abstractions which combine to present a pow-

**Services** Systems boundaries are represented by asynchronous functions called *services*. They provide a symmetric and uniform API: the same abstraction represents both clients and servers.

**Filters** Application-agnostic concerns (e.g. timeouts, retries, authentication) are encapsulated by *filters* which compose to build services from multiple independent modules.

# Error handling in functional way

---

```
module Option =
  let ofChoice choice =
    match choice with
    | Choice1of2 value -> Some value
    | Choice2of2 _ -> None

module AsyncOption =
  type AsyncOption<'a> = Async<'a Option>

  let handler (operation:Async<'a>) : AsyncOption<'a> = ...
```

# Error handling in functional way

---

```
let downloadAsyncImage(blobReference:string) : Async<Image> = async {
    let! container = Helpers.getCloudBlobContainerAsync()
    let blockBlob = container.GetBlockBlobReference(blobReference)
    use memStream = new MemoryStream()
    do! blockBlob.DownloadToStreamAsync(memStream)
    return Bitmap.FromStream(memStream)
}
```

```
downloadAsyncImage "Bugghina001.jpg"
|> AsyncOption.handler
|> Async.map(fun imageOpt ->
    match imageOpt with
    | Some(image) -> image.Save("ImageFolder\Bugghina.jpg")
    | None -> log "There was a problem downloading the image")
|> Async.Start
```

# Preserving Exception semantic - Result Type

---

```
type Result<'TSuccess, 'TFailure> =
| Success of 'TSuccess
| Failure of 'TFailure

module Result =
    let ofChoice value =
        match value with
        | Choice1Of2 value -> Success value
        | Choice2Of2 e -> Failure e

module AsyncResult =
    typeAsyncResult <'a> = Async<Result<'a, exn>>

    let handler (operation:Async<'a>) :AsyncResult<'a> = ...
```

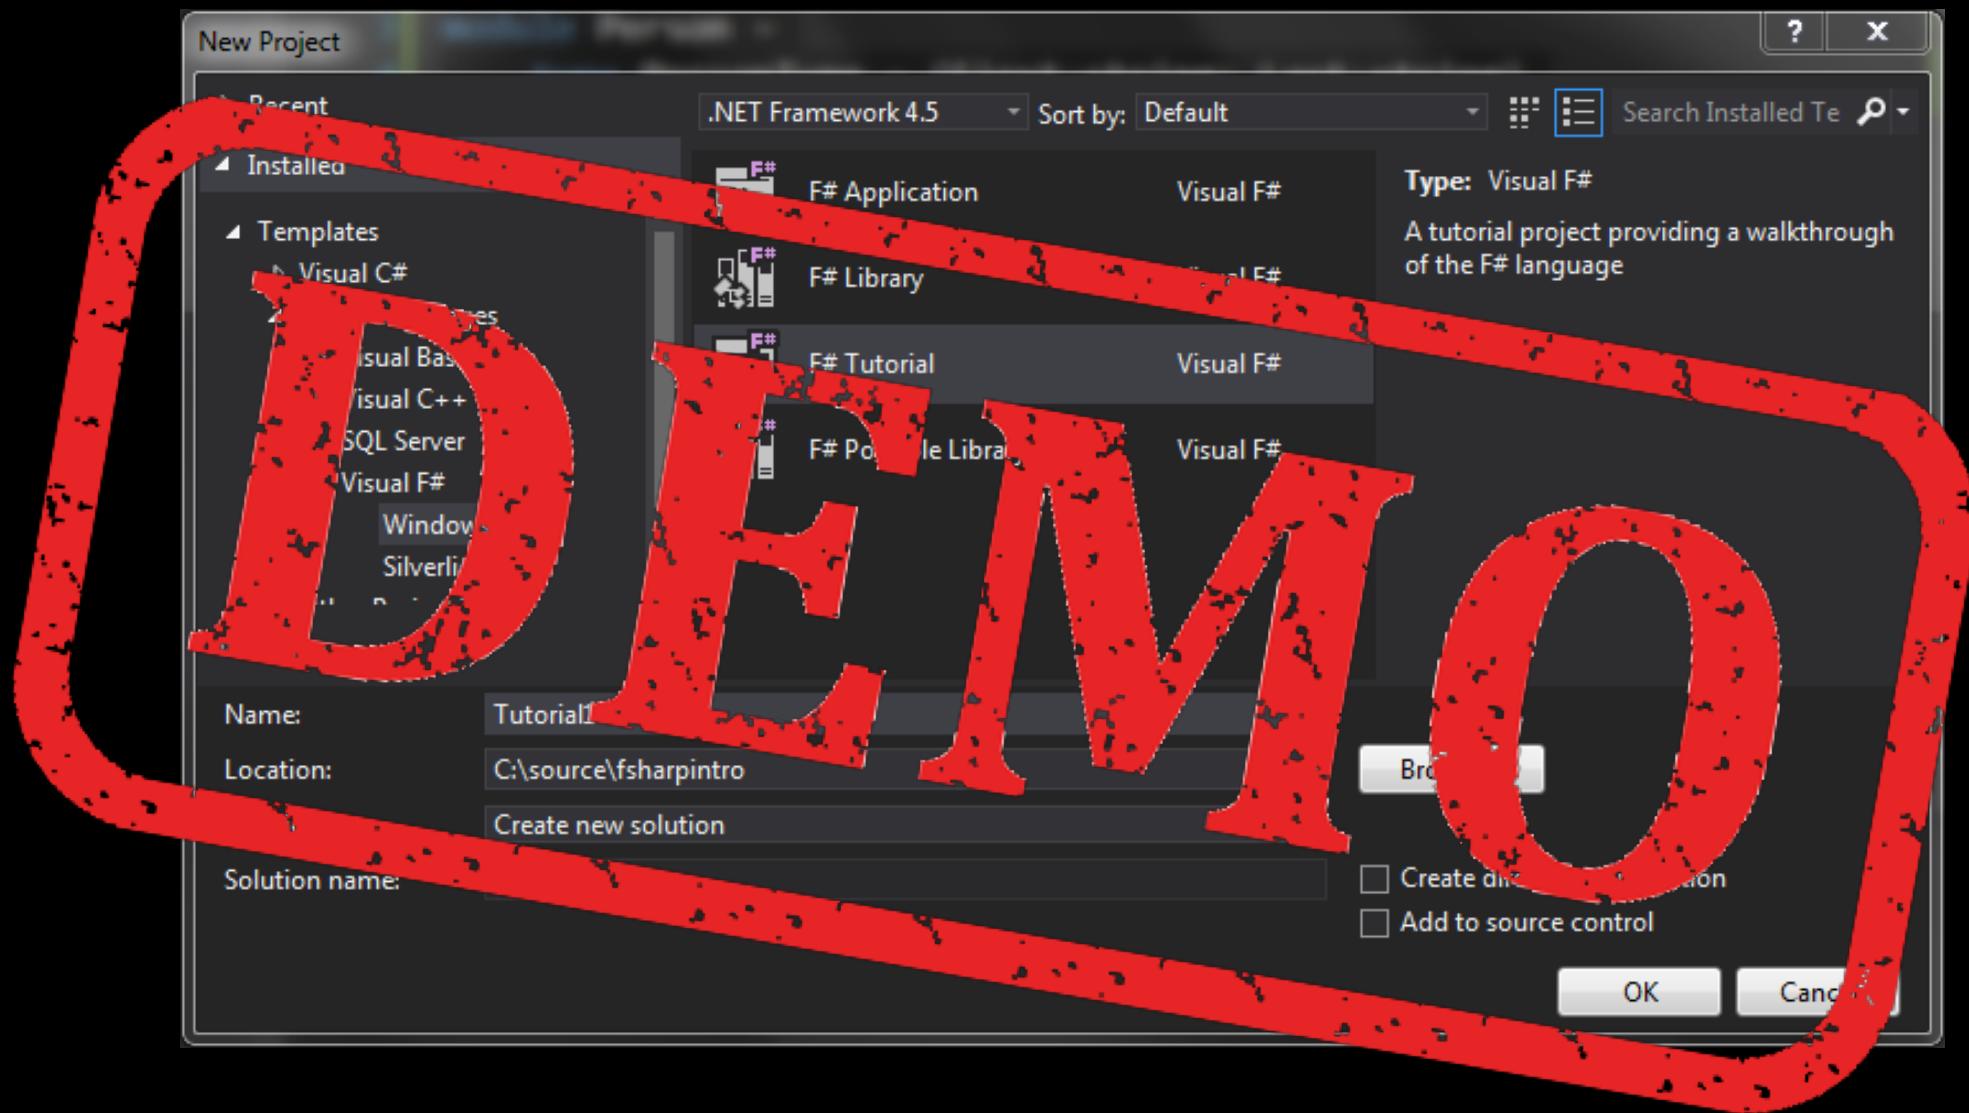
# AsyncResult monadic operators

---

```
module AsyncResult =
  let retn (value:'a) : AsyncResult<'a> =  value |> Ok |> async.Return

  let map (selector : 'a -> 'b) (asyncResult : AsyncResult<'a>) =  ...

  let bind (selector : 'a -> AsyncResult<'b>)
            (asyncResult : AsyncResult<'a>) =  ...
```



# Agent & Reactive Extensions for Extensibility & Interoperability

## **Reactive Extension (Rx)**

- Coordination and composition of event streams
- LINQ-based API

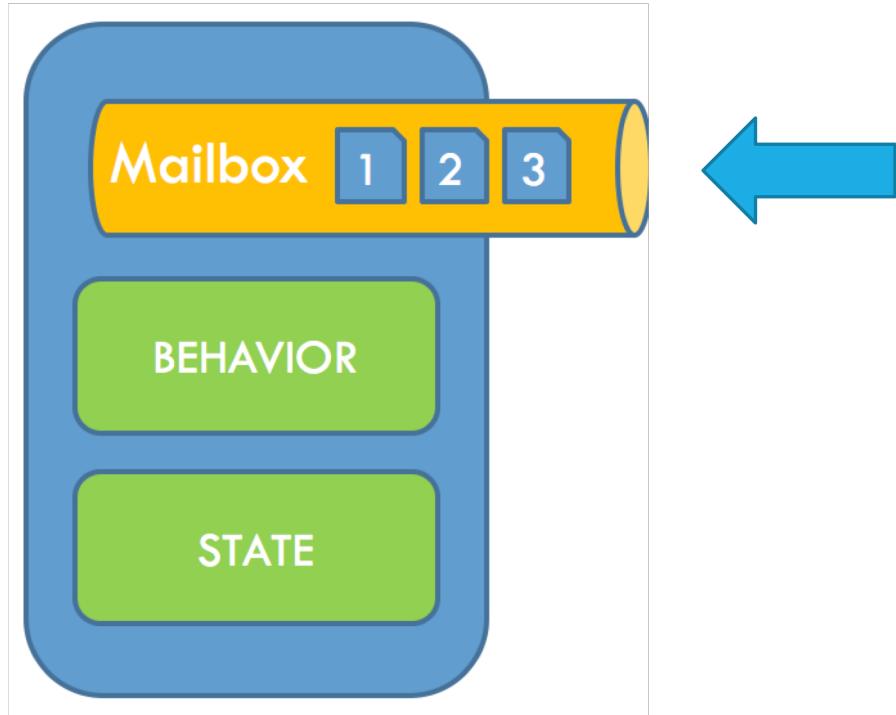
## **MailboxProcessor (MBP)**

- Building blocks for message passing and parallelizing
- Explicit control over how data is buffered and moved

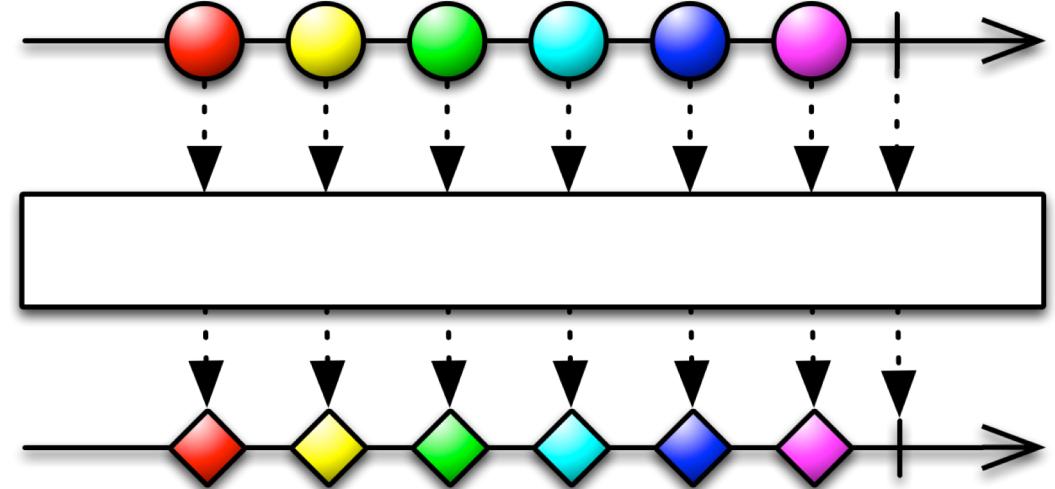
# Streaming complex processing

---

Agent/Actor model



Reactive Extensions



# TPL Dataflow & Reactive Extension

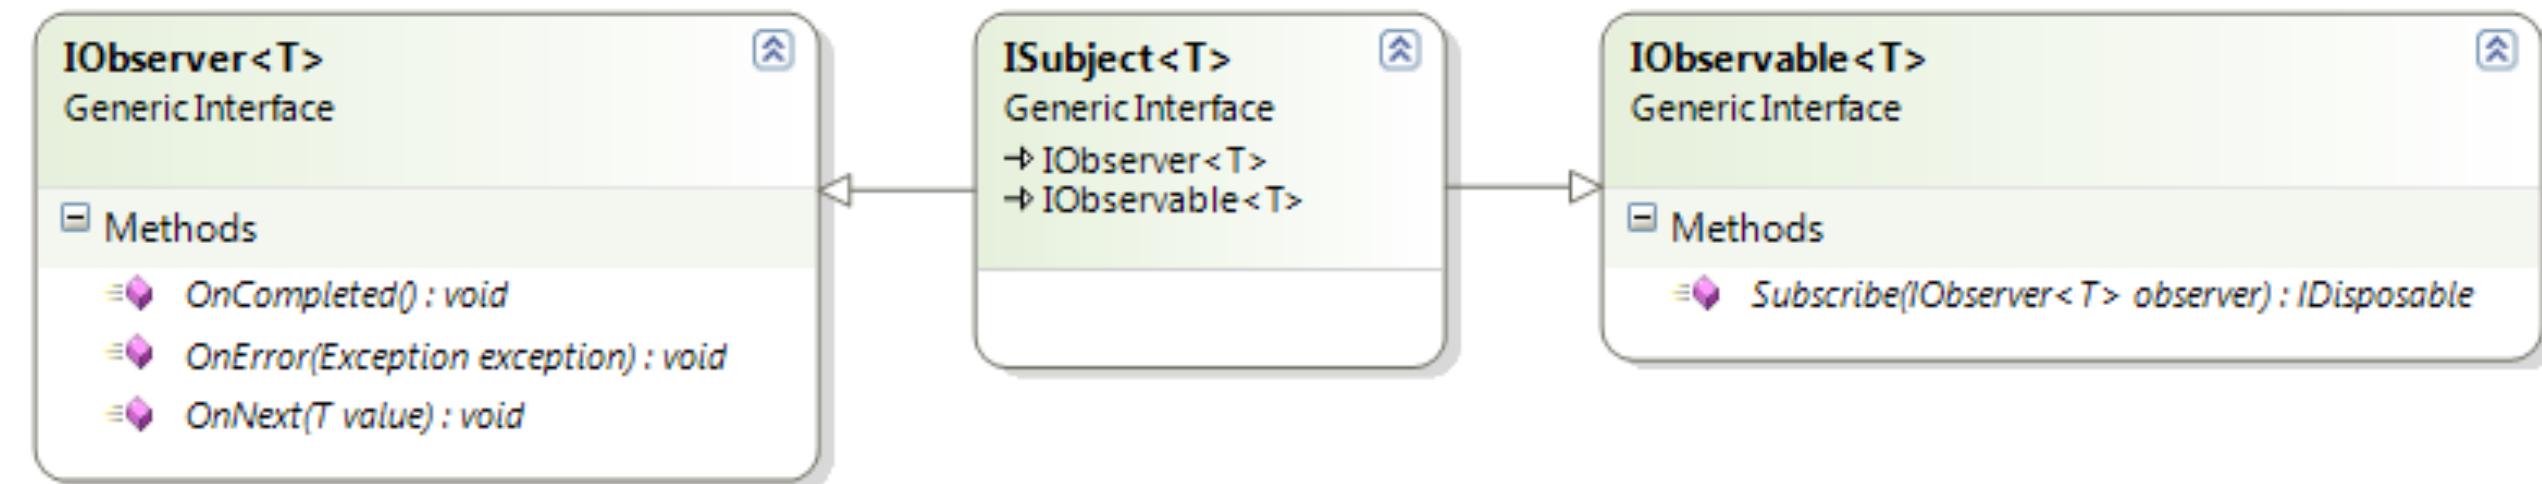
---



Reactive Extensions

# What are Reactive Extensions

Rx is a library for composing asynchronous and event-based programs using observable sequences and LINQ-style query operators



# IObserver & IObservable

---

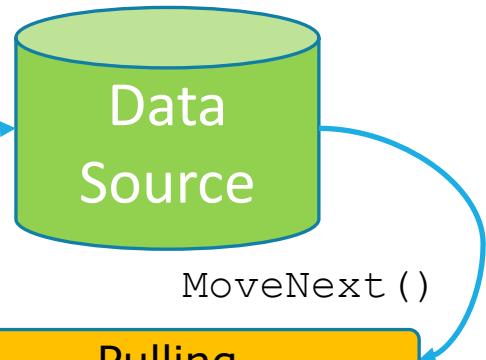
```
type IObserver<'a> = interface
    abstract OnCompleted : unit -> unit
    abstract OnError : exn -> unit
    abstract OnNext : 'a -> unit
end
```



```
type IObservable<'a> = interface
    abstract Subsribe : IObserver<'a> -> IDisposable
end
```

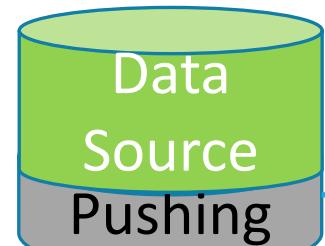
(2) The `IEnumerable`\`IEnumerator` pattern pulls data from source, which blocks the execution if there is no data available

## Interactive



(1) The consumer asks for new data

## Reactive



(2) The `IObservable`\`IObserver` pattern receives a notification from the source when new data is available, which is pushed to the consumer

(1) The source notifies the consumer that new data is available

# Fundamental Abstractions

## Adapting the observer pattern

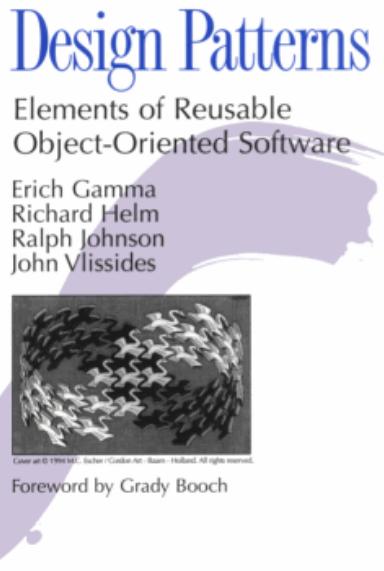
- Ensuring duality with the enumerator pattern
- More compositional approach

```
interface IObservable<out T>
{
    IDisposable Subscribe(IObserver<T> observer);
}
```

```
interface IObserver<in T>
{
    void OnNext(T value);
    void OnError(Exception error);
    void OnCompleted();
}
```

### Notification grammar

OnNext\* (OnError | OnCompleted)?



# Stream processing basics – Combinators

---

```
let observable = Observable  
    .Interval(TimeSpan.FromSeconds(1.0))  
  
observable.Subscribe (printfn "%d")
```

# Stream processing basics – Combinators

---

```
let observable = Observable  
    .Interval(TimeSpan.FromSeconds(1.0))  
    .GroupBy(fun x -> x % 3L)
```

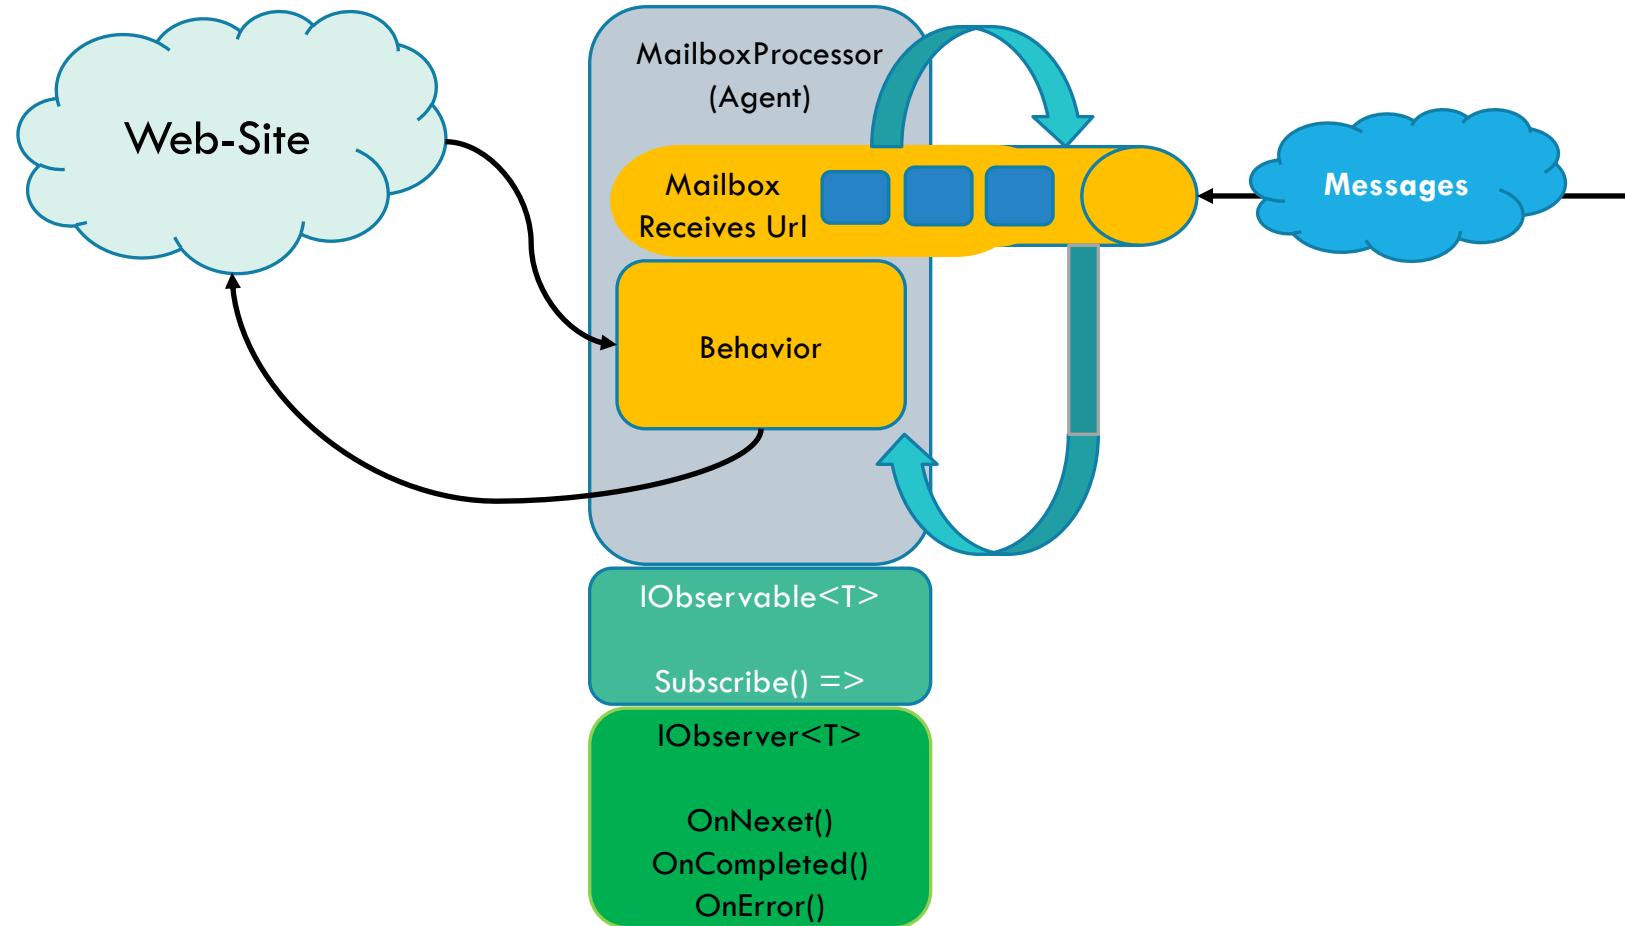
```
let printKeyValue key value =  
    printfn "%s %d" (String('*', key)) value
```

```
observable.Subscribe(fun obs ->  
    obs.Subscribe (printKeyValue (int obs.Key)) |> ignore)
```

# Reactive Agents



# Agent as Observable



# MailboxProcessor and Rx

---

```
inputBuffer.LinkTo(compressor, linkOptions);
compressor.LinkTo(encryptor, linkOptions);

encryptor.AsObservable()

    .Scan(new Dictionary<int, EncryptDetails>(), 0),
    (state, msg) => Observable.FromAsync(async() => {
        Dictionary<int,EncryptDetails> details, int lastIndexProc) = state;
        details.Add(msg.Sequence, msg);

        return (details, lastIndexProc);
    }) .SingleAsync())
    .SubscribeOn(TaskPoolScheduler.Default).Subscribe();
```

Lab :

Reactive WebCrawler with Agents

# Summary

---

- » The free lunch is over
- » Applications must be built for concurrency
- » Shared state concurrency is hard
- » Agent can be compose and parallelize
- » Alternative concurrency models
  - Message passing concurrency

# Measure twice and cut once

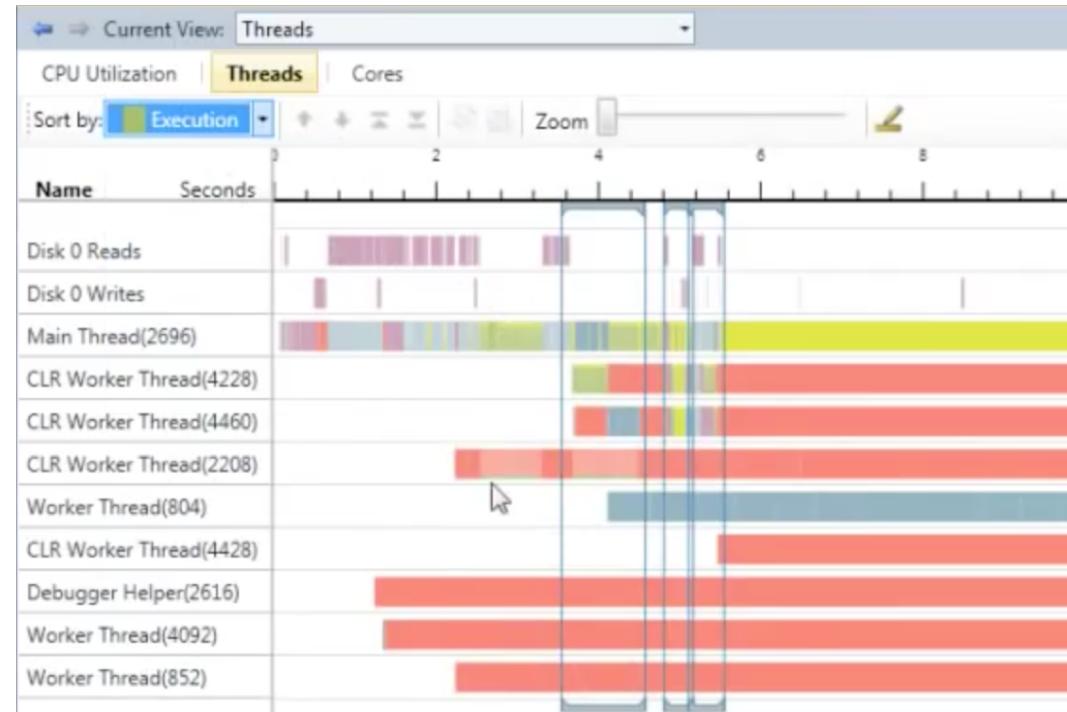
---



# Assume nothing – Profile everything

---

- Be scientific
- Do test multiple implementations
- Don't set out to confirm your bias
- Instrument and profile your code





*The tools we use have a profound (and devious!) influence on our thinking habits, and, therefore, on our thinking abilities.*

-- Edsger Dijkstra

<https://ti.to/luteceo-workshops/concurrent-functional-programming-in-net/en>



**LUTECEO**  
workshops

Concurrent Functional Programming in .NET

📅 April 25th–26th, 2019     🗺 Paris, France

A photograph of Riccardo Terrell, a man with glasses and a dark sweater, speaking at a conference. He is gesturing with his right hand. To his left is a blue circular button with the word "Courses". Below the photo is a QR code and the "skills matter" logo.

**Functional  
Concurrency in  
.NET with C# & F#  
with Riccardo  
Terrell**

A two day course to master the necessary practices to build concurrent and scalable programs in .NET

Visit [skillsmatter.com/courses/589-functional-concurrency-in-dotnet-with-csharp-and-fsharp-with-riccardo-terrell](https://skillsmatter.com/courses/589-functional-concurrency-in-dotnet-with-csharp-and-fsharp-with-riccardo-terrell) for more information

**29th - 30th April 2019 in London**

<https://skillsmatter.com/courses/589-functional-concurrency-in-dotnet-with-csharp-and-fsharp-with-riccardo-terrell>